
Producing Minimal Hook Placements to Enforce Authorization Policies

Divya Muthukumaran, Nirupama Talele, and Trent Jaeger

Penn State University

With Vinod Ganapathy (Rutgers) and Gang Tan (Lehigh)


Security is not a 'blocker'

"It isn't clear this qualifies as a blocker under any circumstances. The importance of security increases only as we are into serious deployment and start becoming a target. First things, first...."

- <https://dev.laptop.org/ticket/260>

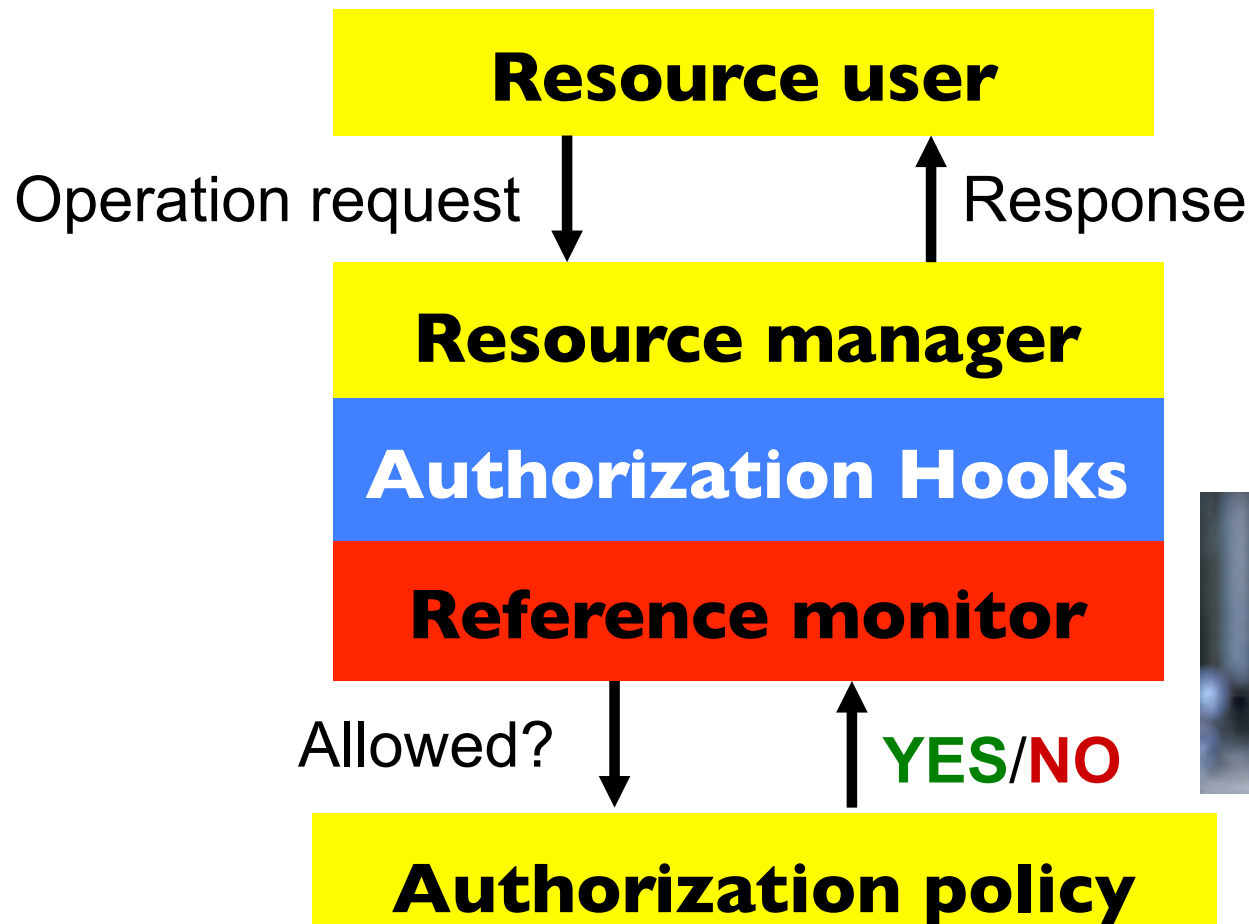
Need a way to protect against applications sniffing each other's keystrokes, which X permits by default.

Need for retroactive security




- Large codebases need retroactive security features.
- X Server, postgres, OpenSSH, Linux Kernel, etc.
- Different problems:
 - Privilege separation, Memory errors, *Authorization*, Logging.

What is authorization ?




⟨Alice, /etc/passwd, *File_Read*⟩

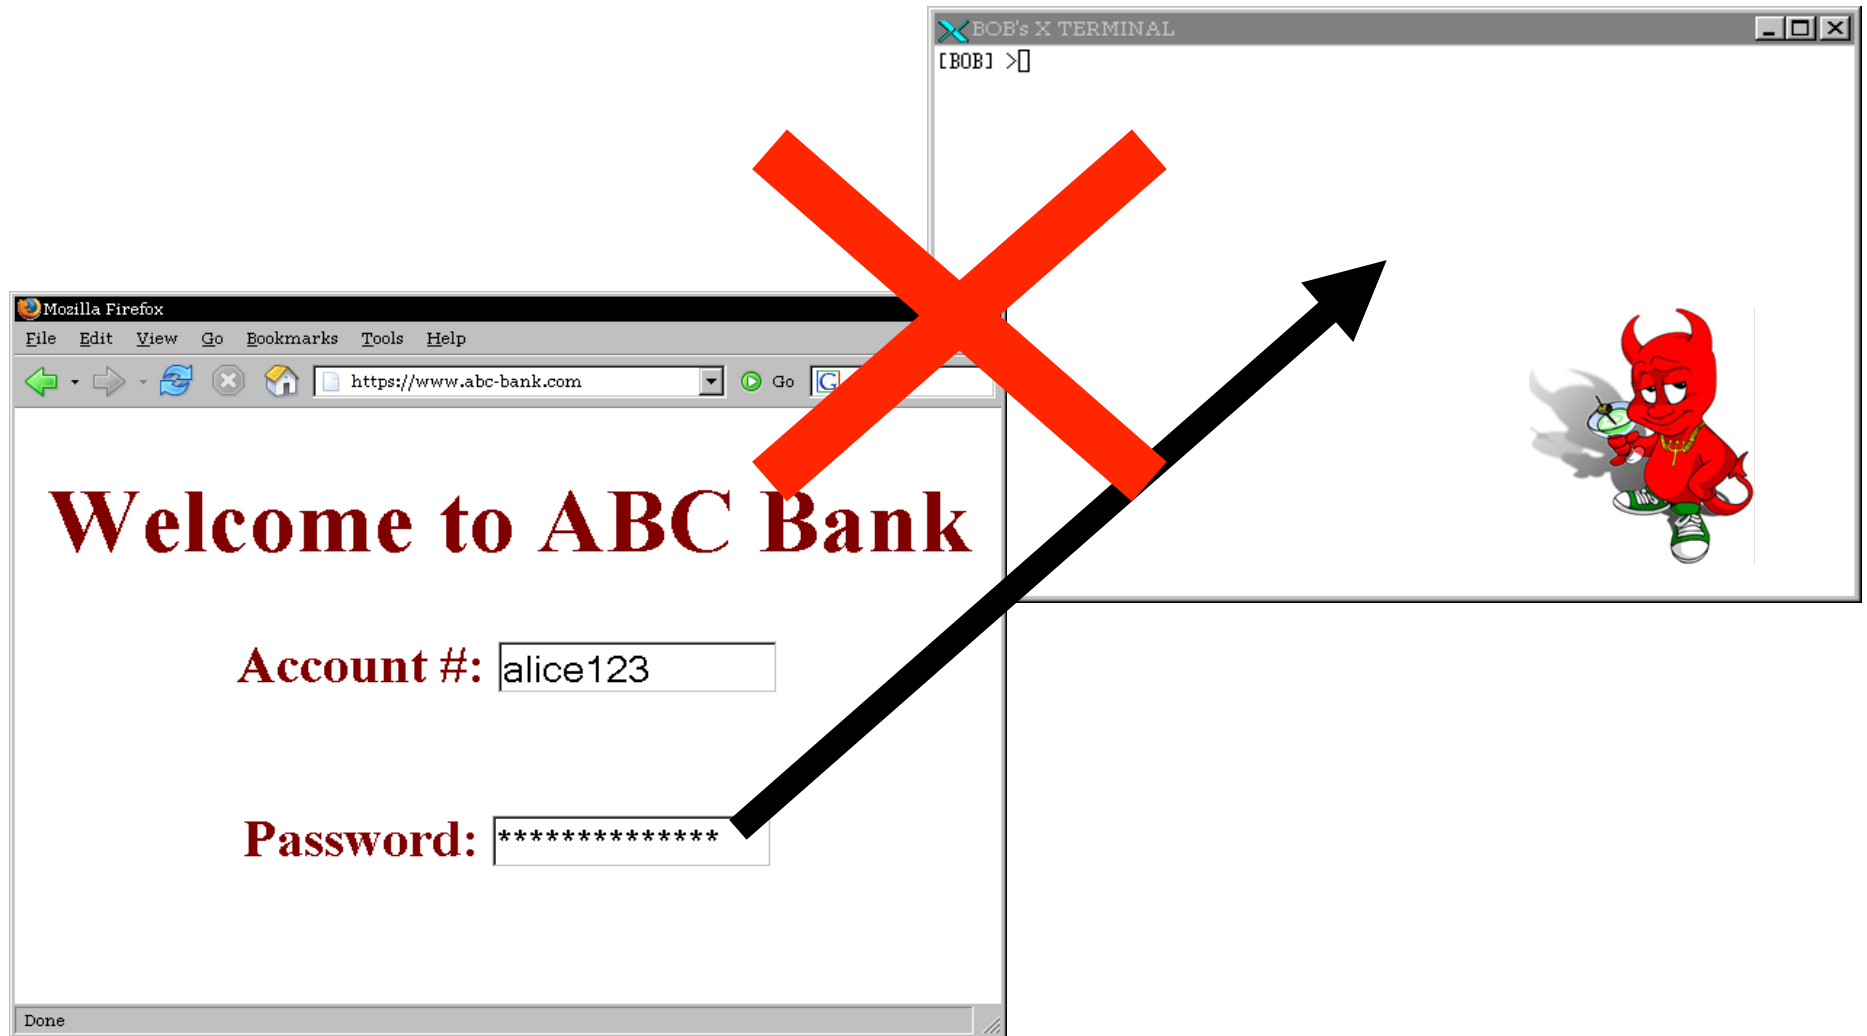
X Server



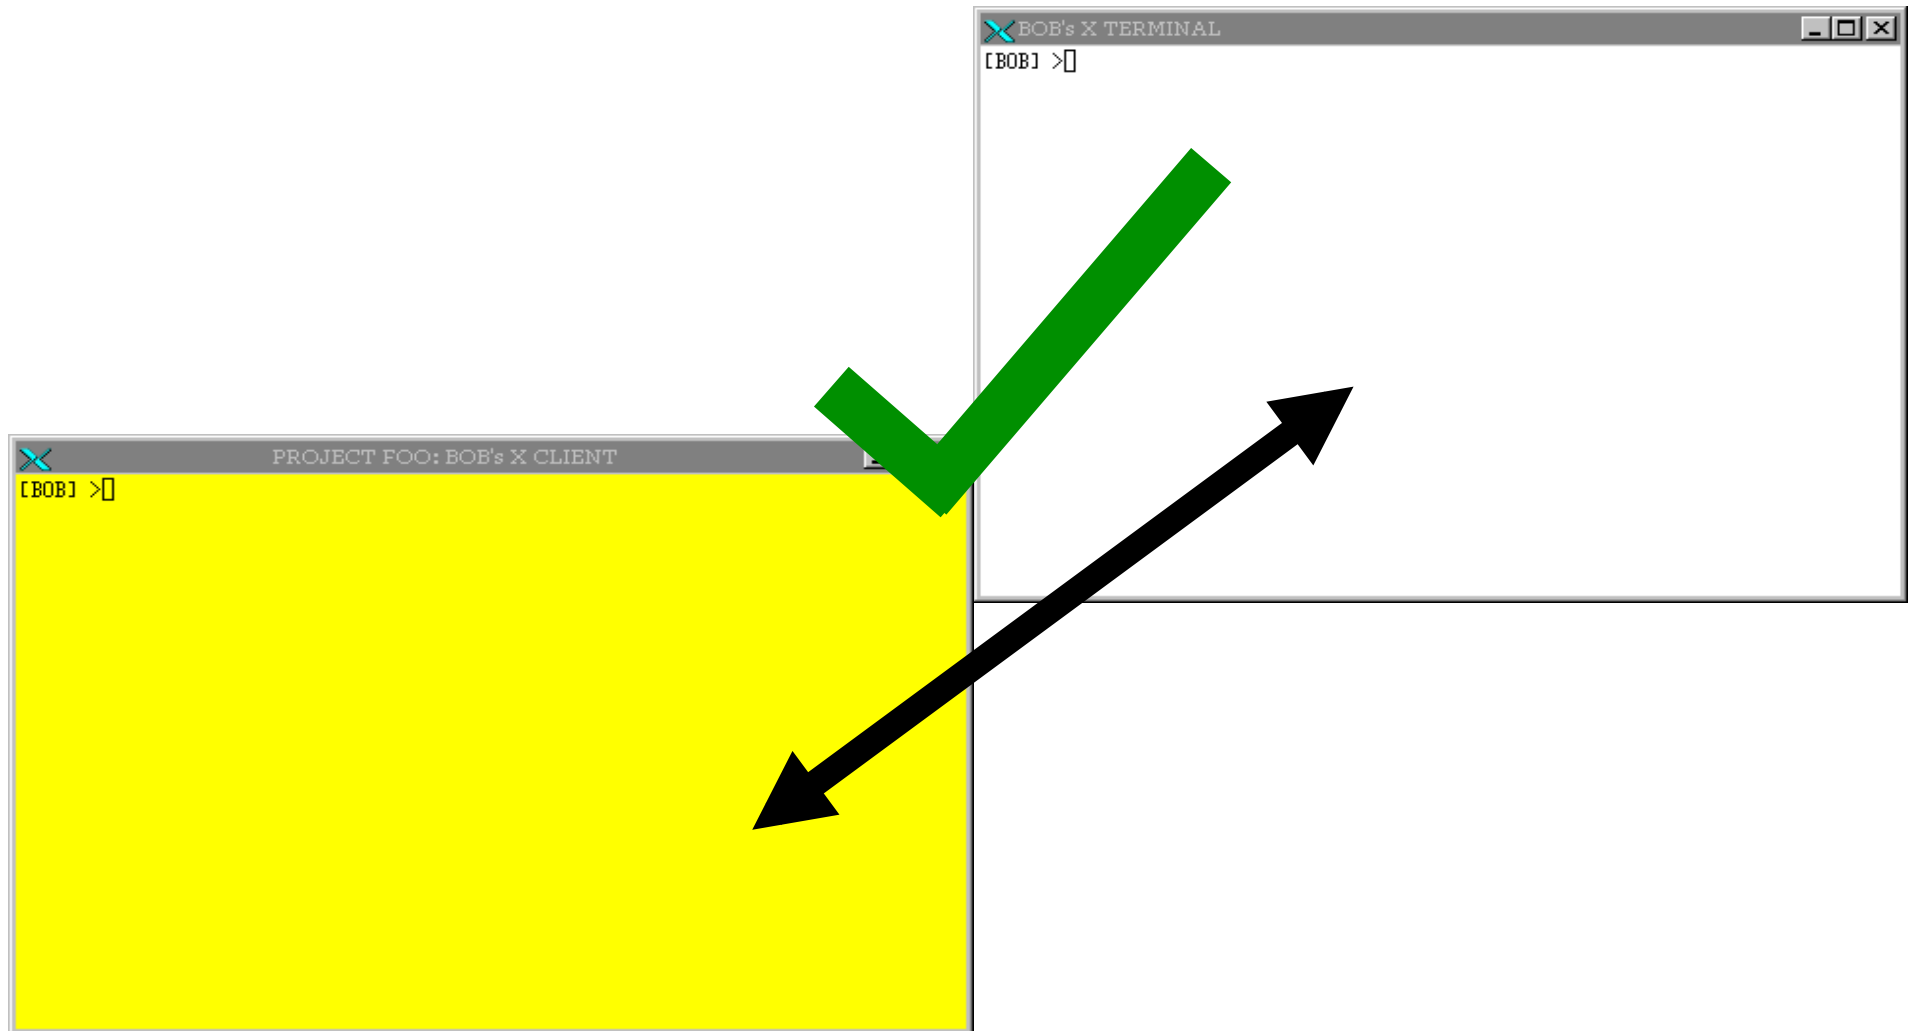
X Server



X Server




X Server



Retrofitting code for authorization

Need systematic techniques to retrofit legacy code for security



Painstaking, manual procedure



- ▣ X11 ~ proposed 2003, upstreamed 2007, changing to date. [Kilpatrick et al., '03]
- ▣ Linux Security Modules ~ 2 years [Wright et al., '02]

*At this point, SE-PostgreSQL has taken up a ***lot* of community resources**, not to mention an **enormous and doubtless frustrating amount of *the lead developer's* time and effort**, thus far **without a single committed patch, or even a consensus as to what it should (or could) do**. Rather than continuing to blunder into the future, I think we need to do a reality check*
- <http://archives.postgresql.org/message-id/20090718160600.GE5172@fetter.org>

Past efforts

Verifying Hook Consistency

Assumptions: Mature Code

Inputs: Existing hook placement.

Examples:

For Kernels [Zhang et al., 2002, Edwards et al., 2002, Tan et al., 2008]

For Web Applications [Sun et al., 2011, RoleCast 2011, FixMeUp 2012]

Placing Hooks Automatically

Assumptions: No hooks

Inputs: Sensitive data types, hook code.

Examples:

Server Applications
[Ganapathy et al., 2005, 2006, 2007]

Thesis Statement

We can **automatically** retrofit server applications with **minimal** number of **authorization hooks** necessary to enforce the desired **authorization policies**.

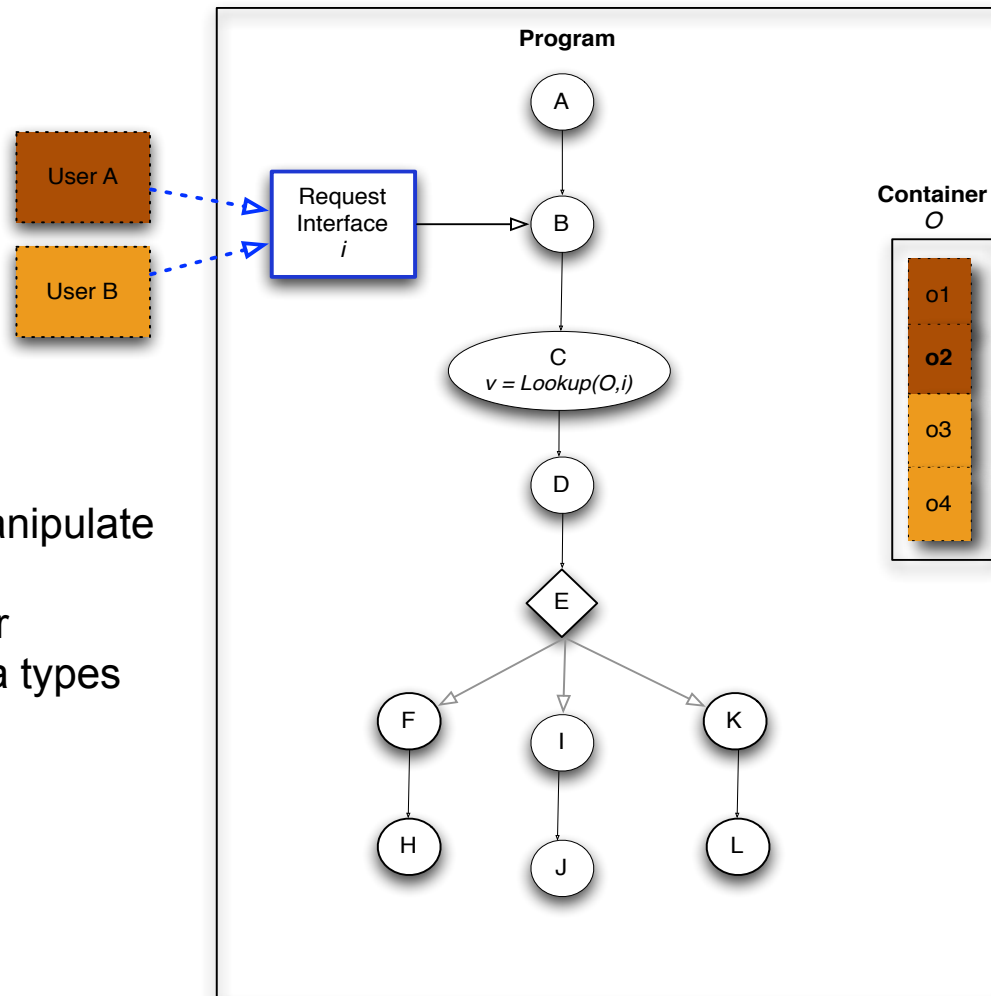
Challenges

- ▣ *Infer security-sensitive **resources** and **operations**.*
- ▣ *Produce **hook placements**.*
 - ▣ Infer locations where hooks need to be placed.
 - ▣ Do not actually place the hooks.

Authorization Tuple

- ▣ [Subject, Resource, Operation]:
- ▣ **Subject**: Entities served by the program. They make requests.
 - ▣ Determined at runtime.
- ▣ **Resources (objects)**: Entities to which access need to be controlled.
 - ▣ Given a large program, how can we identify resources that need access control?
- ▣ **Operations**: Determines what needs to be done with the resources.
 - ▣ Operation is a set of statements. Which sets of statements are security sensitive?

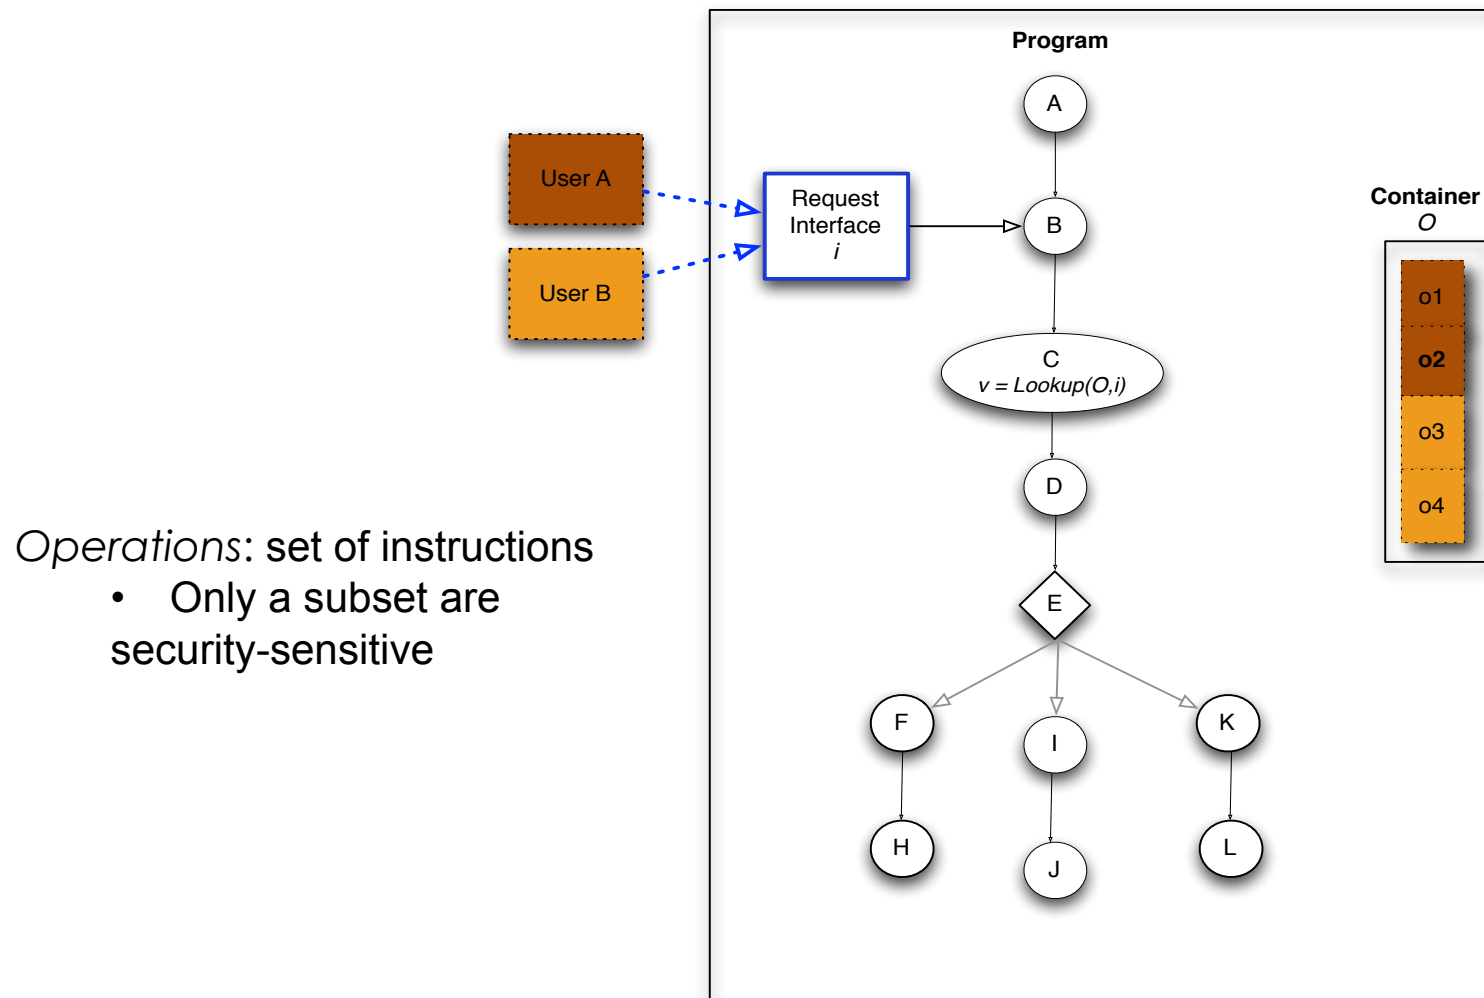
Inferring security-sensitive resources



Resources: Programs manipulate many variables

- 7800 in X Server
- Of over 400 data types

Inferring security-sensitive operations



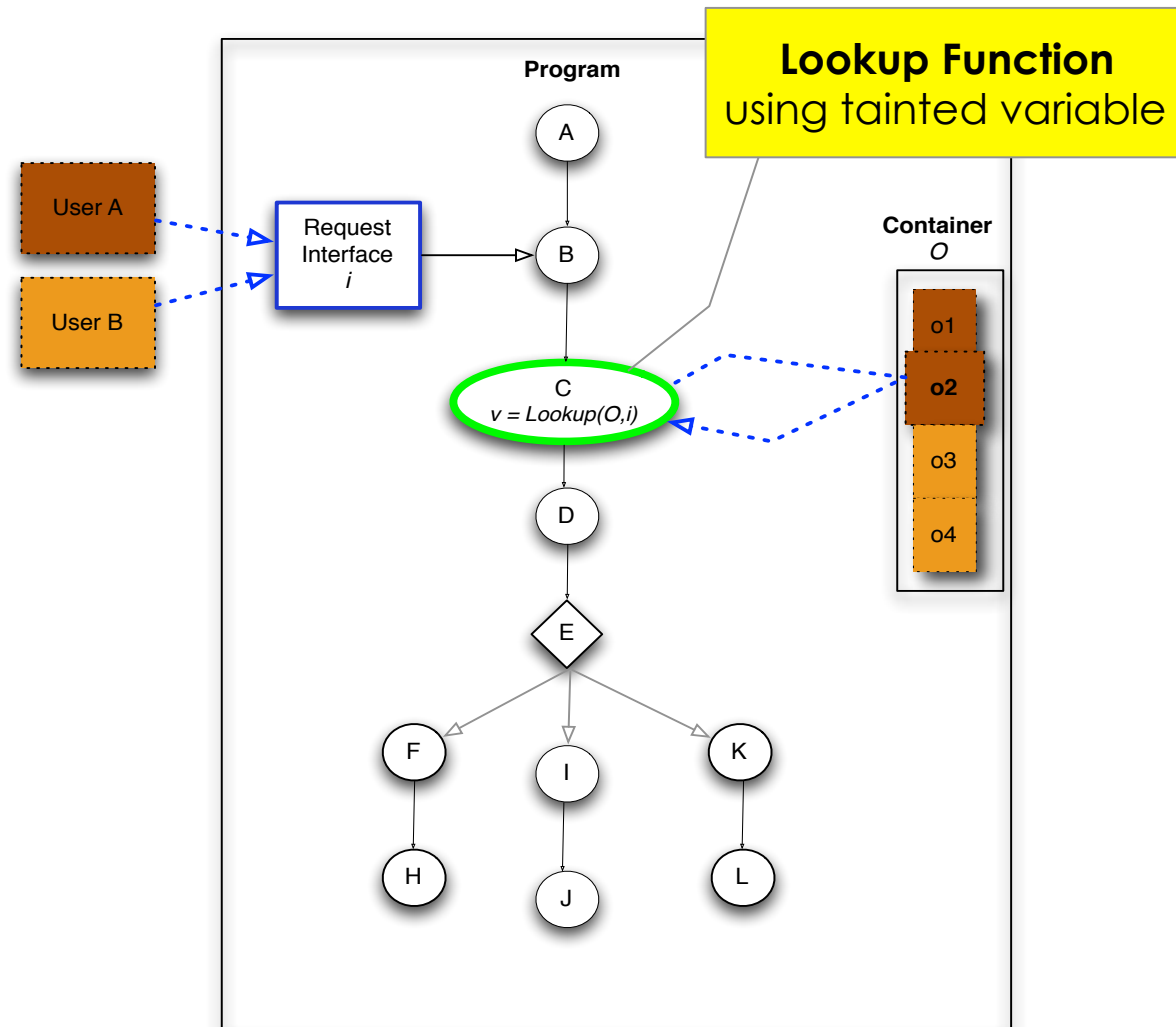
Solution

Requests make choices

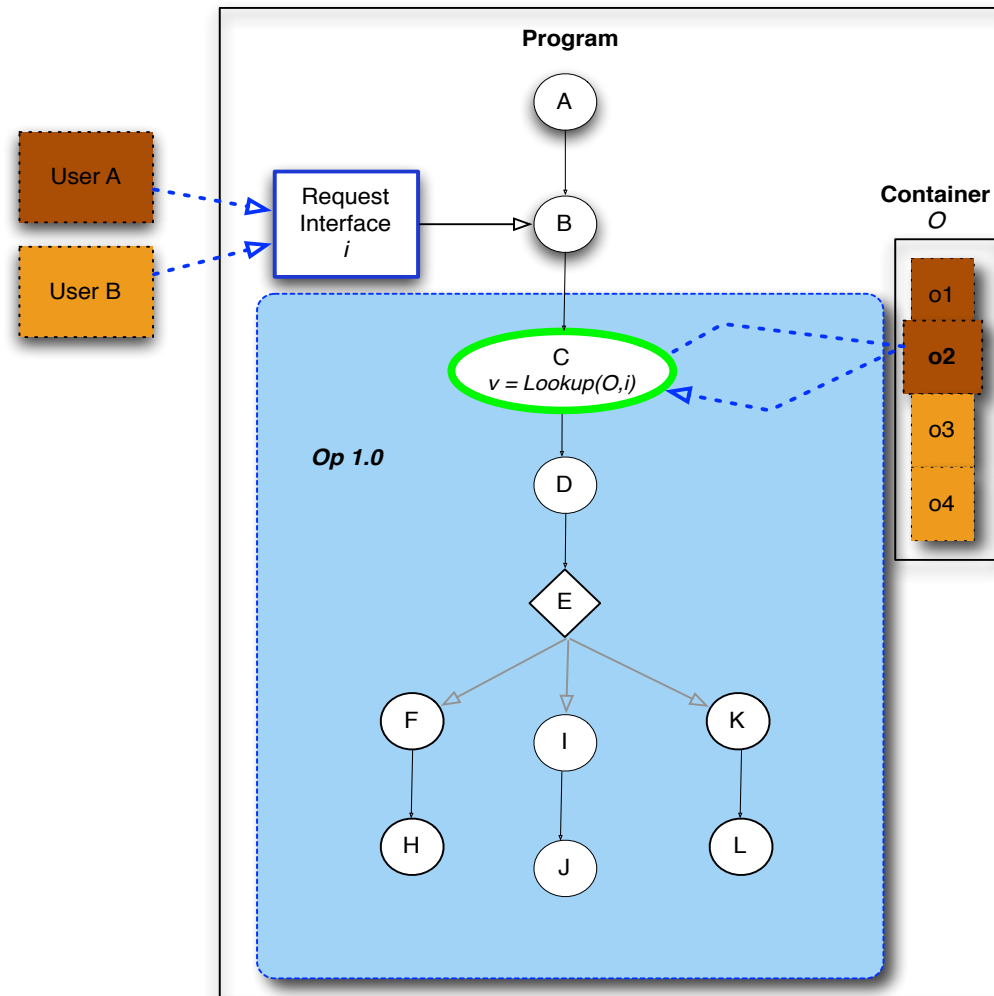
- In servers, *client-request* determines *choices* that client subjects can make in the program
- “Choice”:
 - **Resources:** Determine which *resources* are chosen from containers.
 - **Operations:** Determine which *program path* is selected for execution.



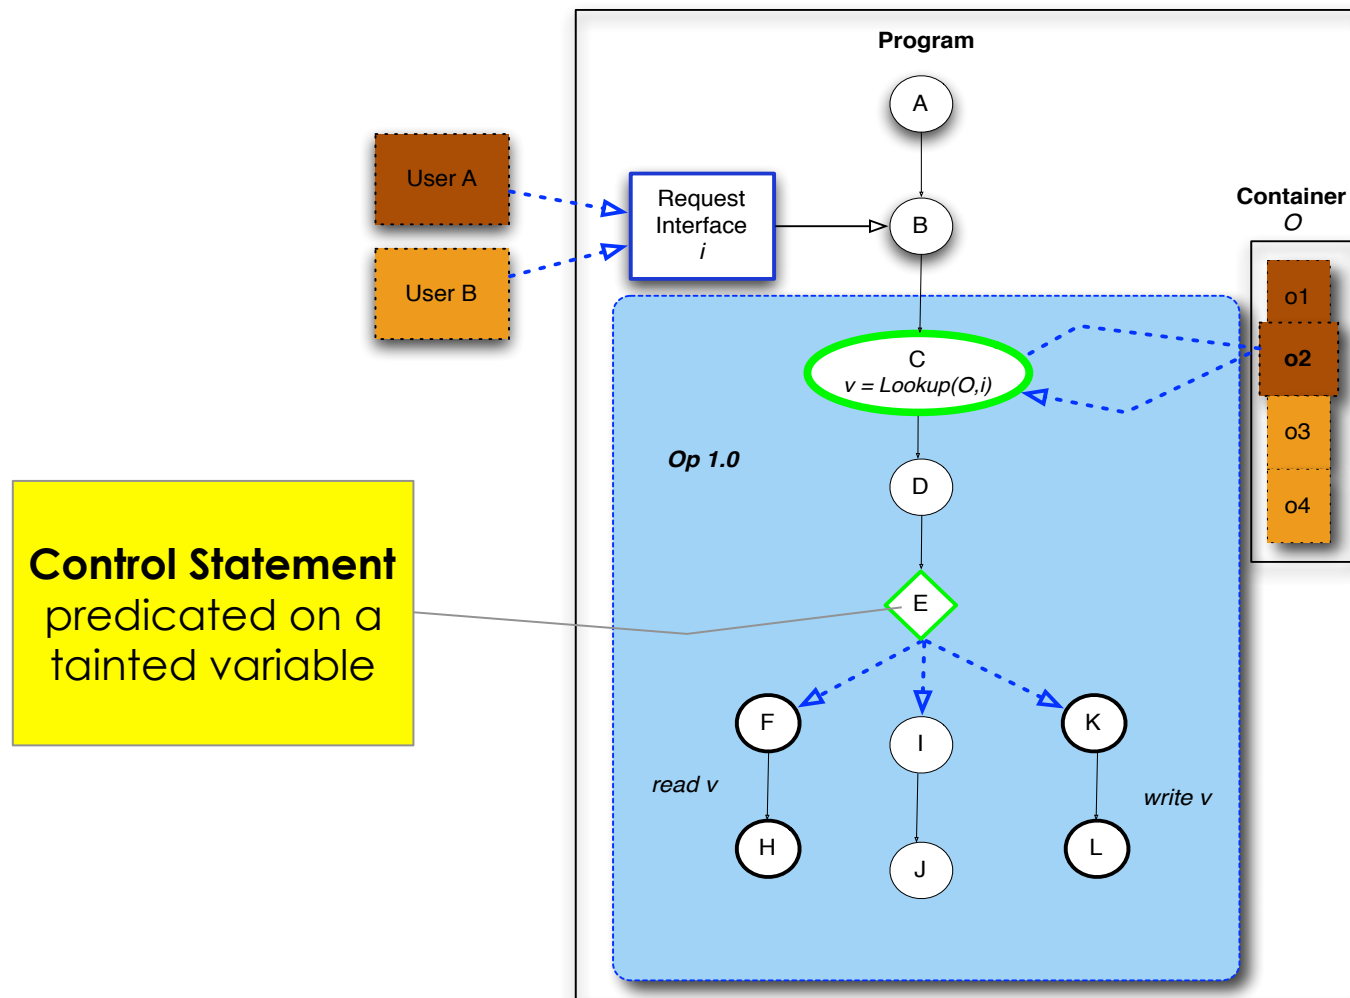
Inferring security-sensitive resources



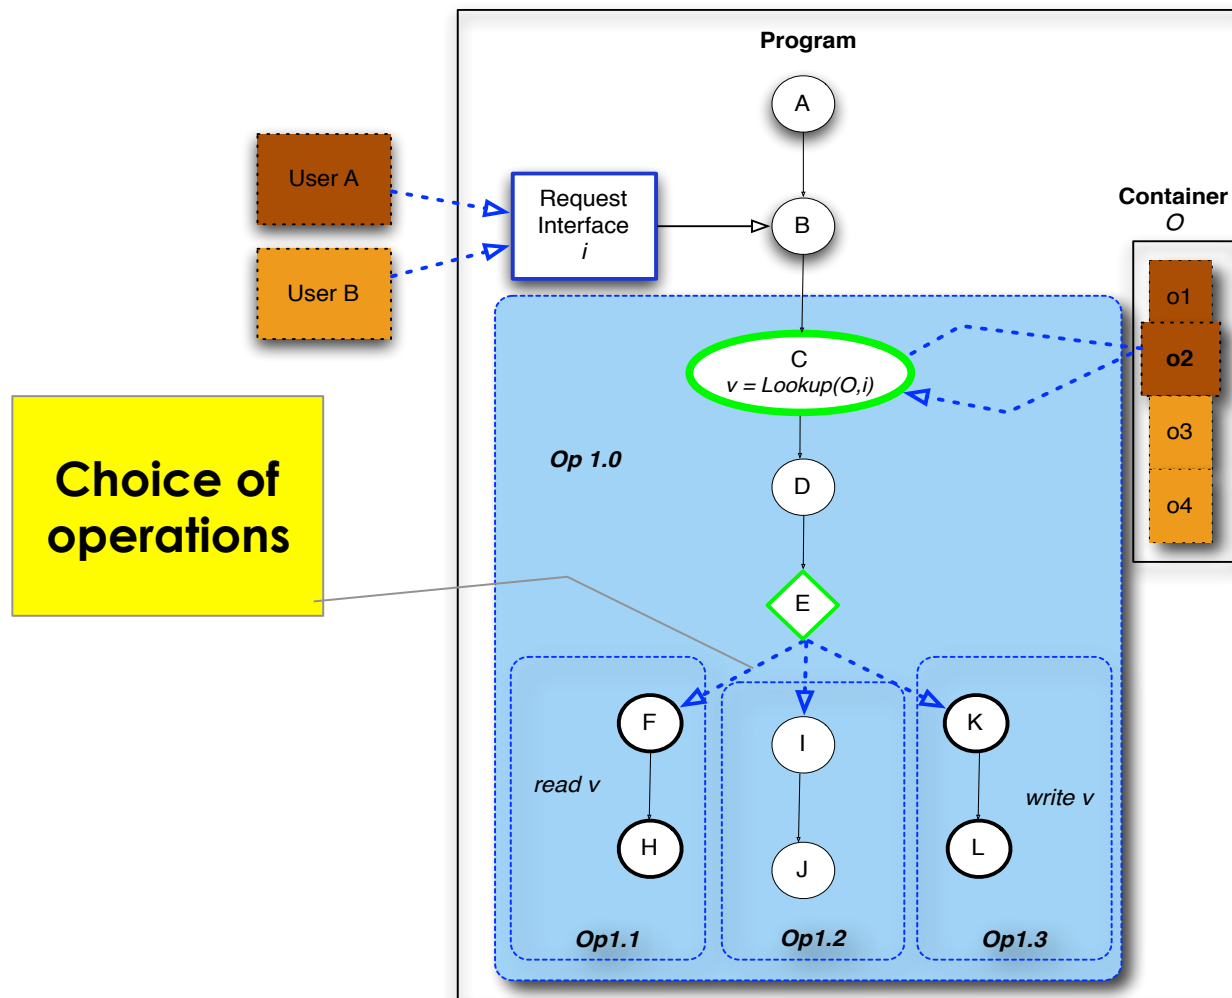
Inferring security-sensitive operations



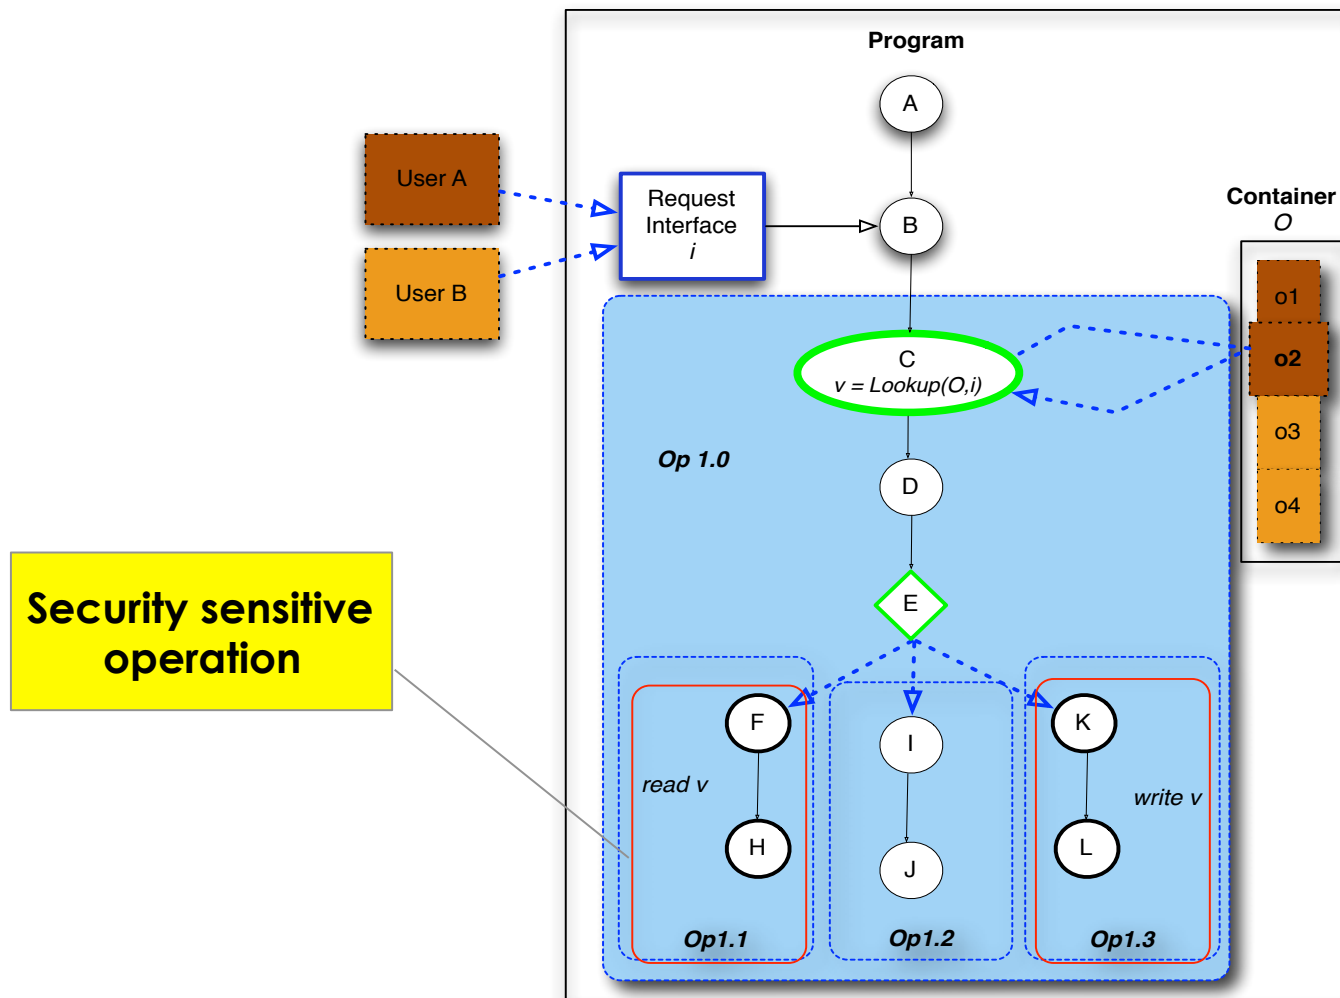
Inferring security-sensitive operations



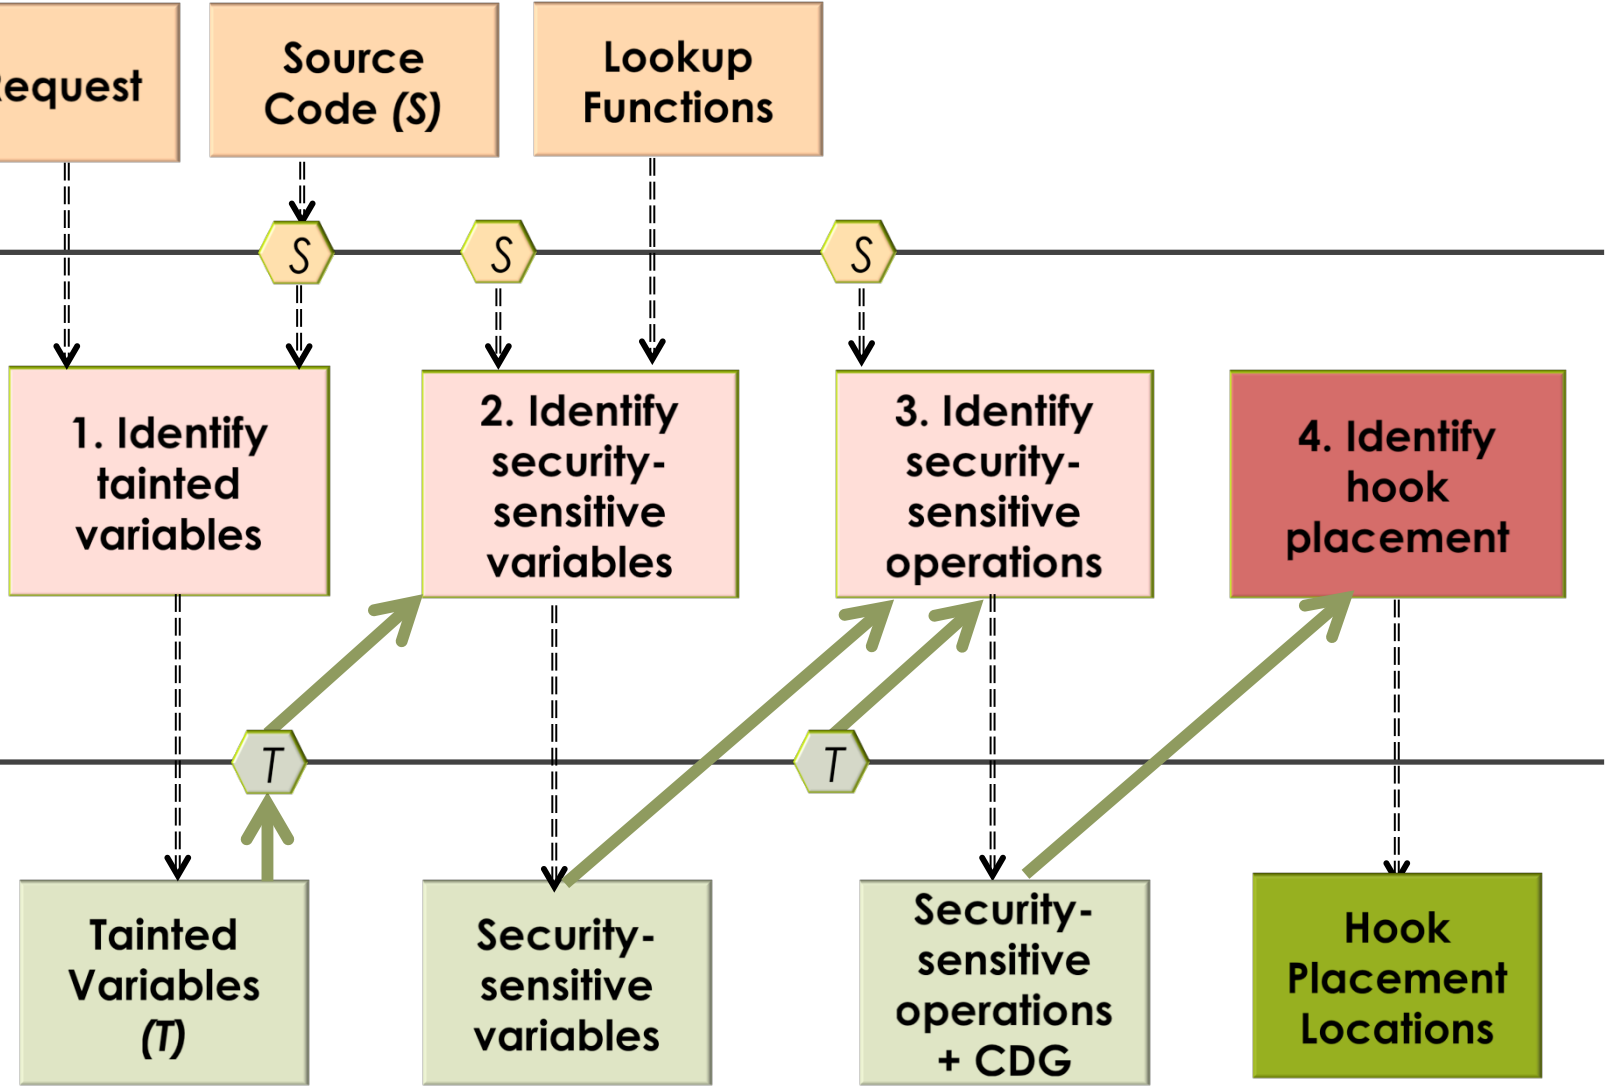
Inferring security-sensitive operations



Inferring security-sensitive operations




Design




Techniques Used

- *Static taint analysis*
 - Identify variables tainted by user request.
 - Identify security-sensitive objects.
- *Control dependence analysis*
 - Identify security-sensitive operations.
 - Hoist and remove redundant hooks.

Control Dependence




Control Flow Graph
(CFG)




Control
Dependence
Graph (CDG)

Control Dependence

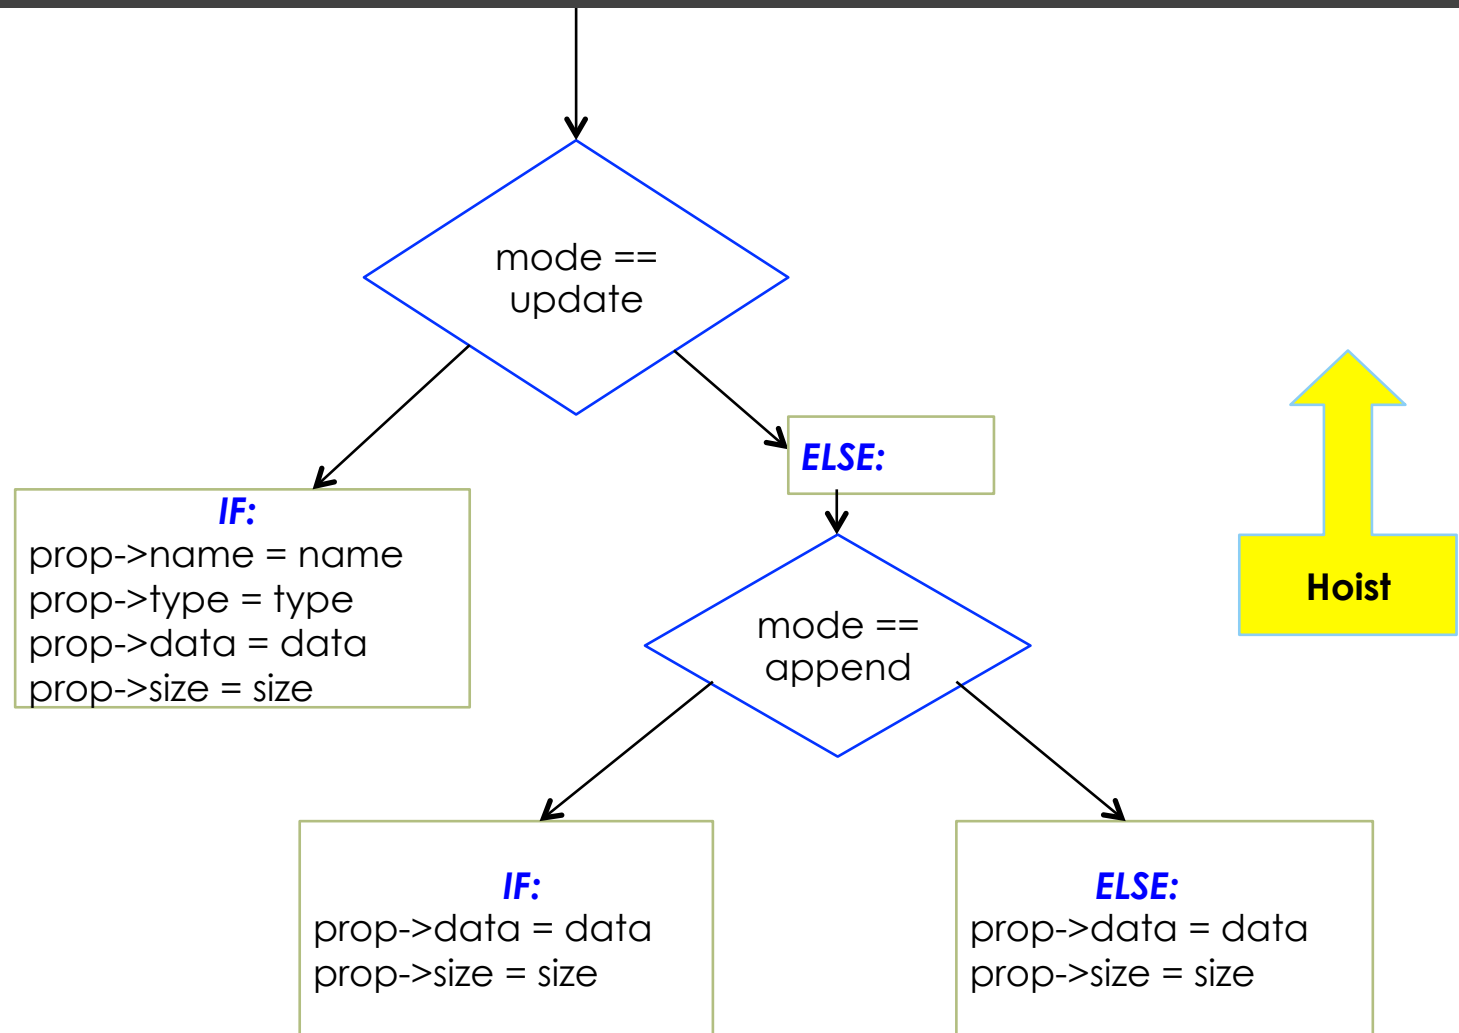


Control Flow Graph
(CFG)

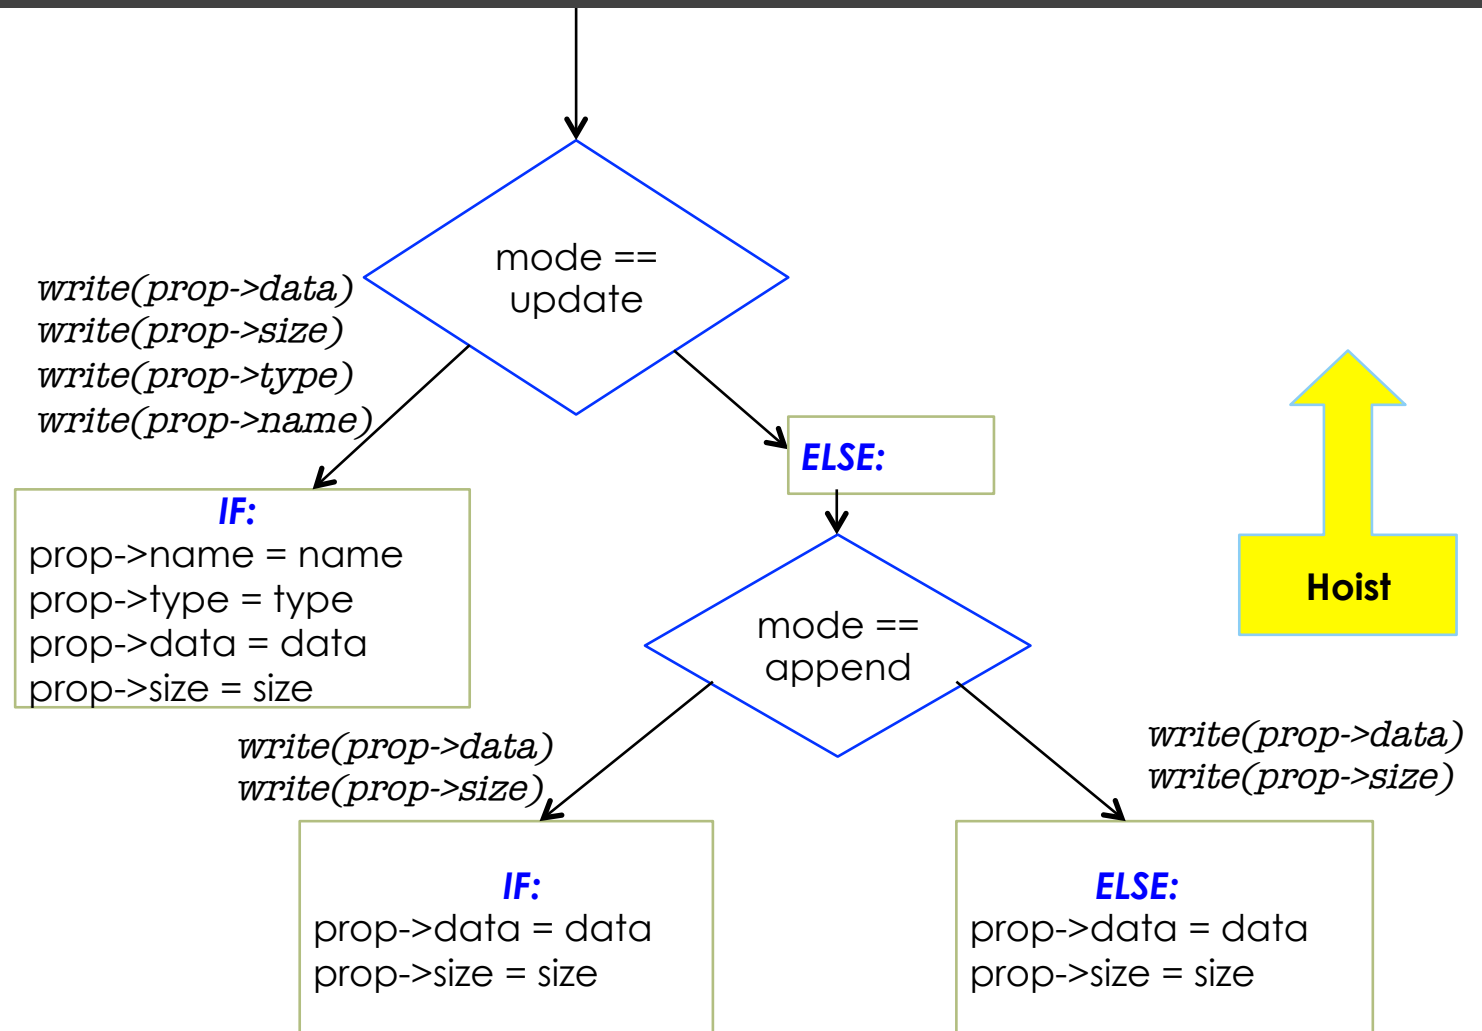


Control
Dependence
Graph (CDG)

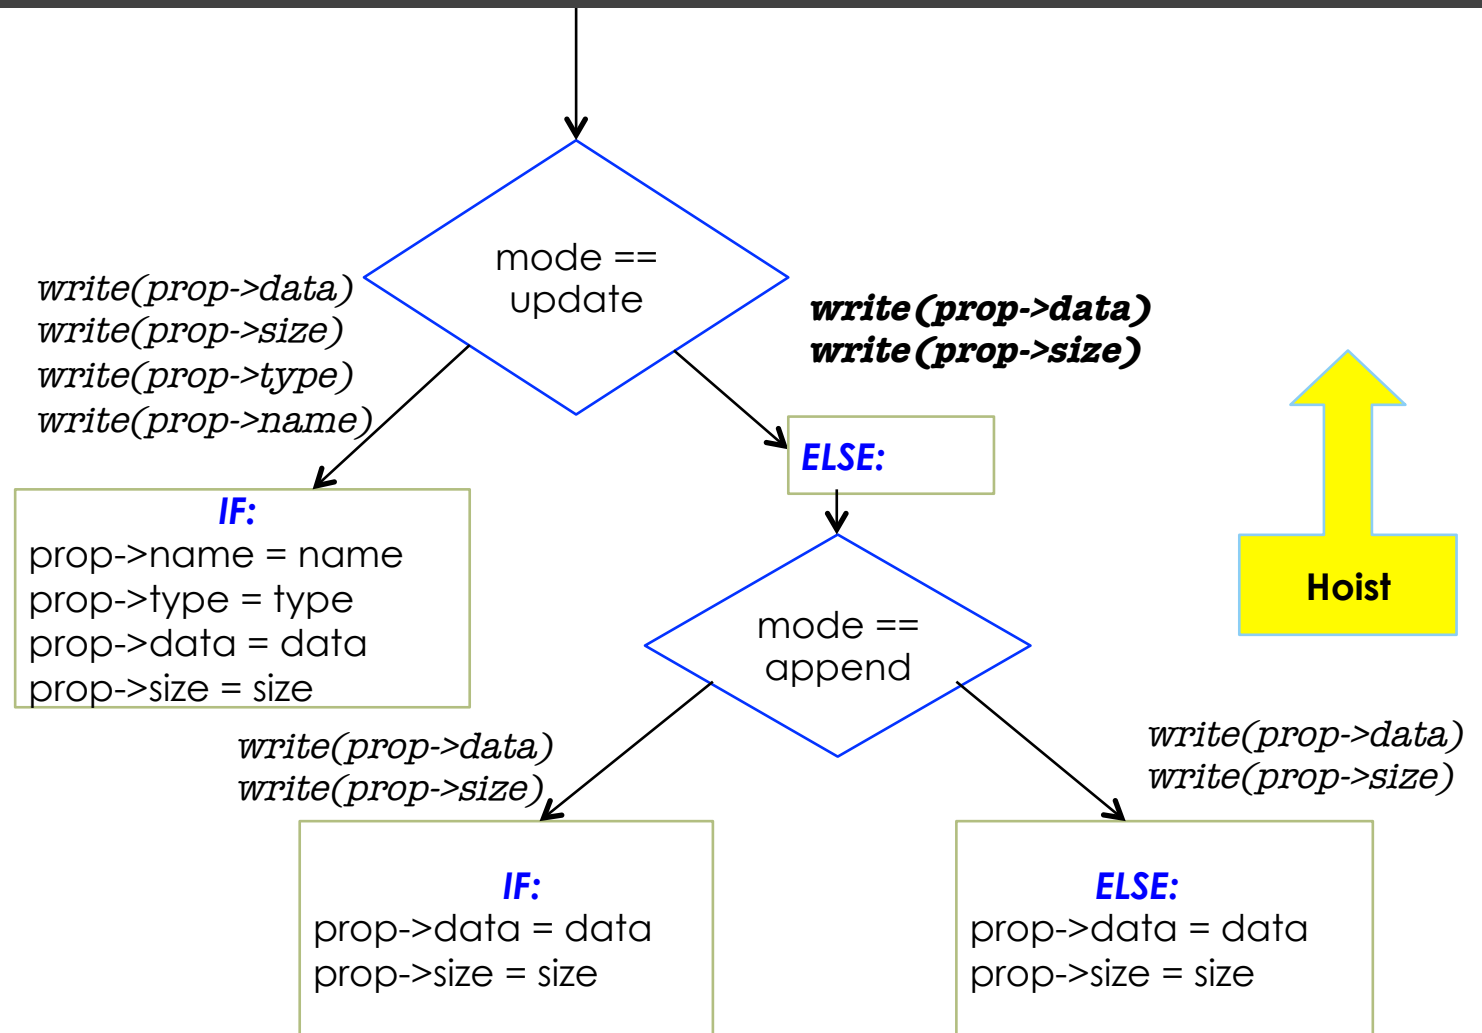
Non-redundant Hook Placement



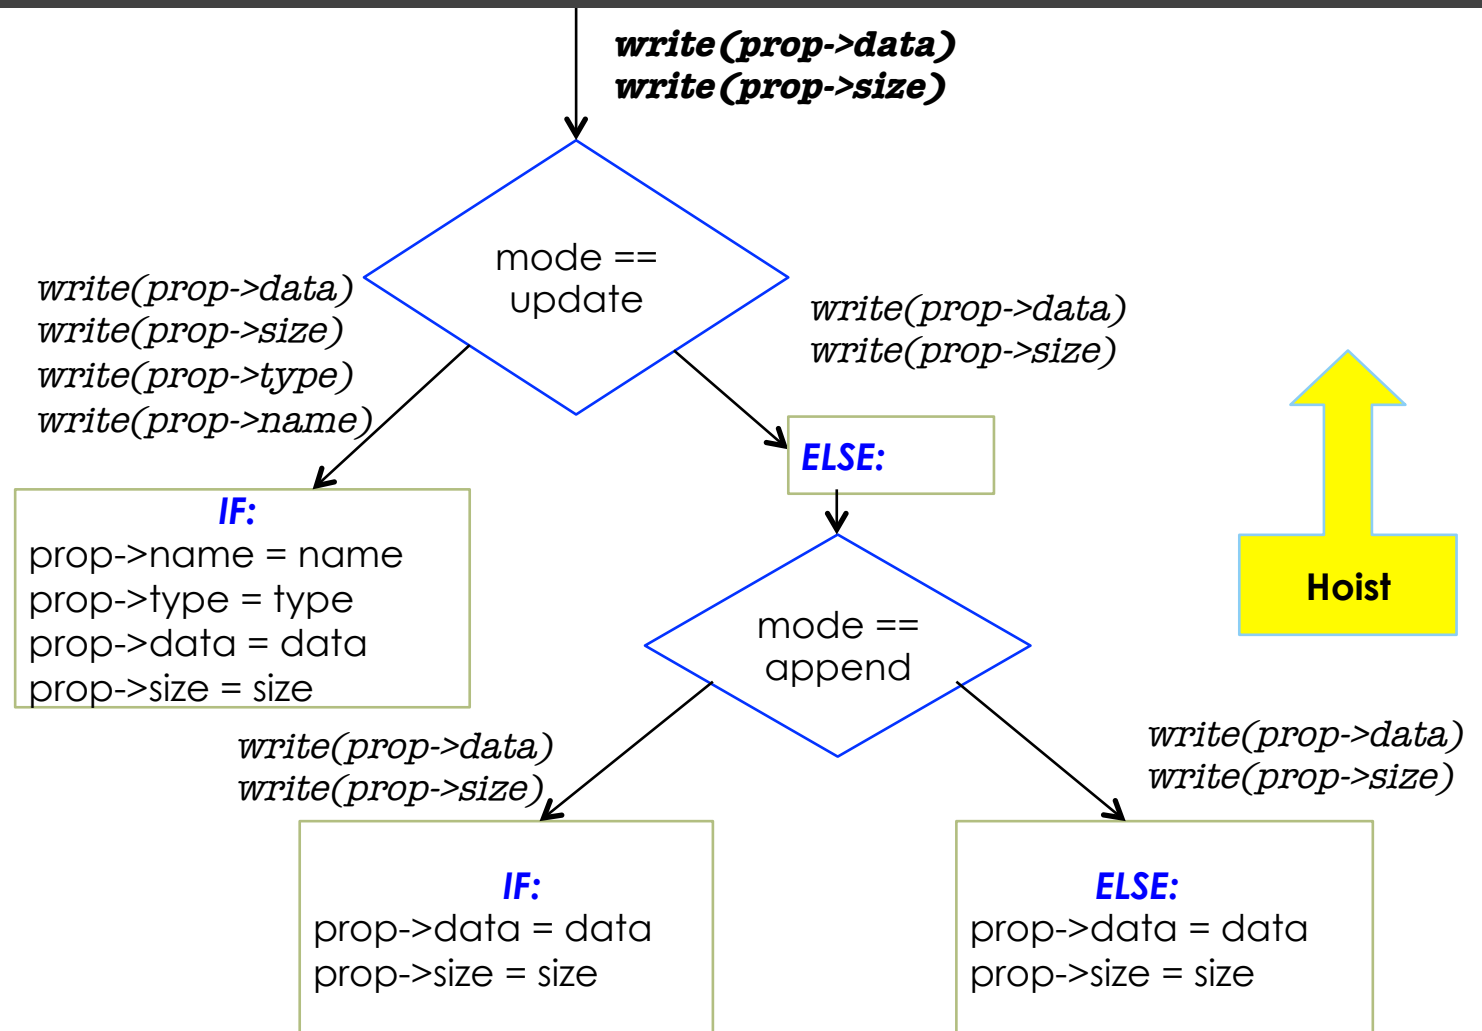
Non-redundant Hook Placement



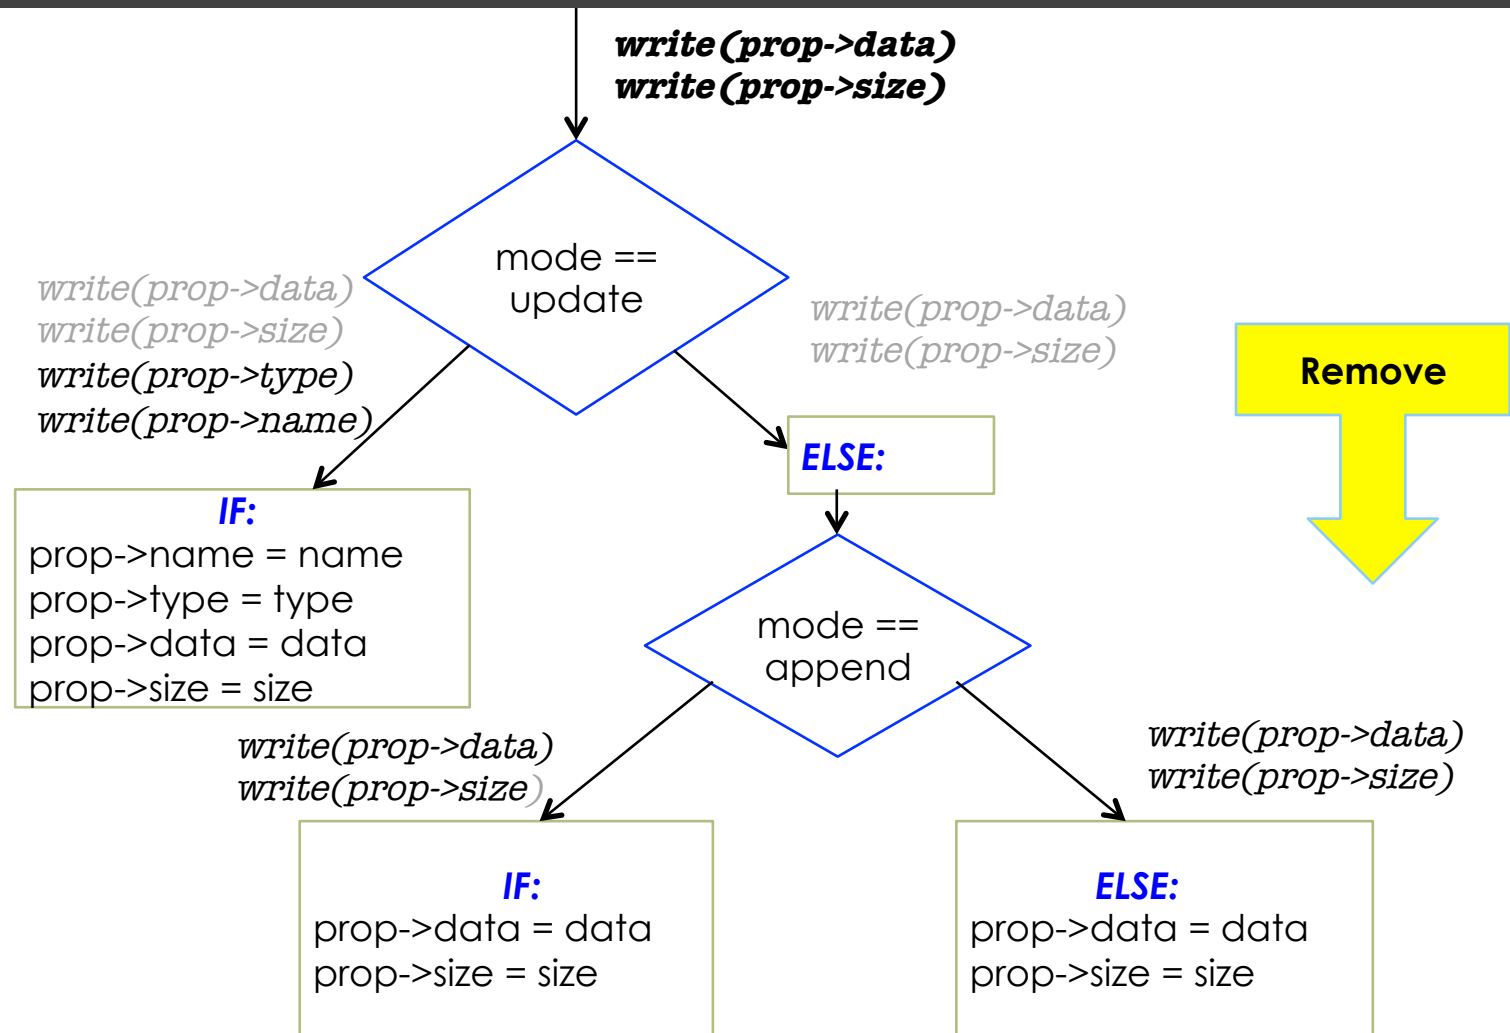
Non-redundant Hook Placement



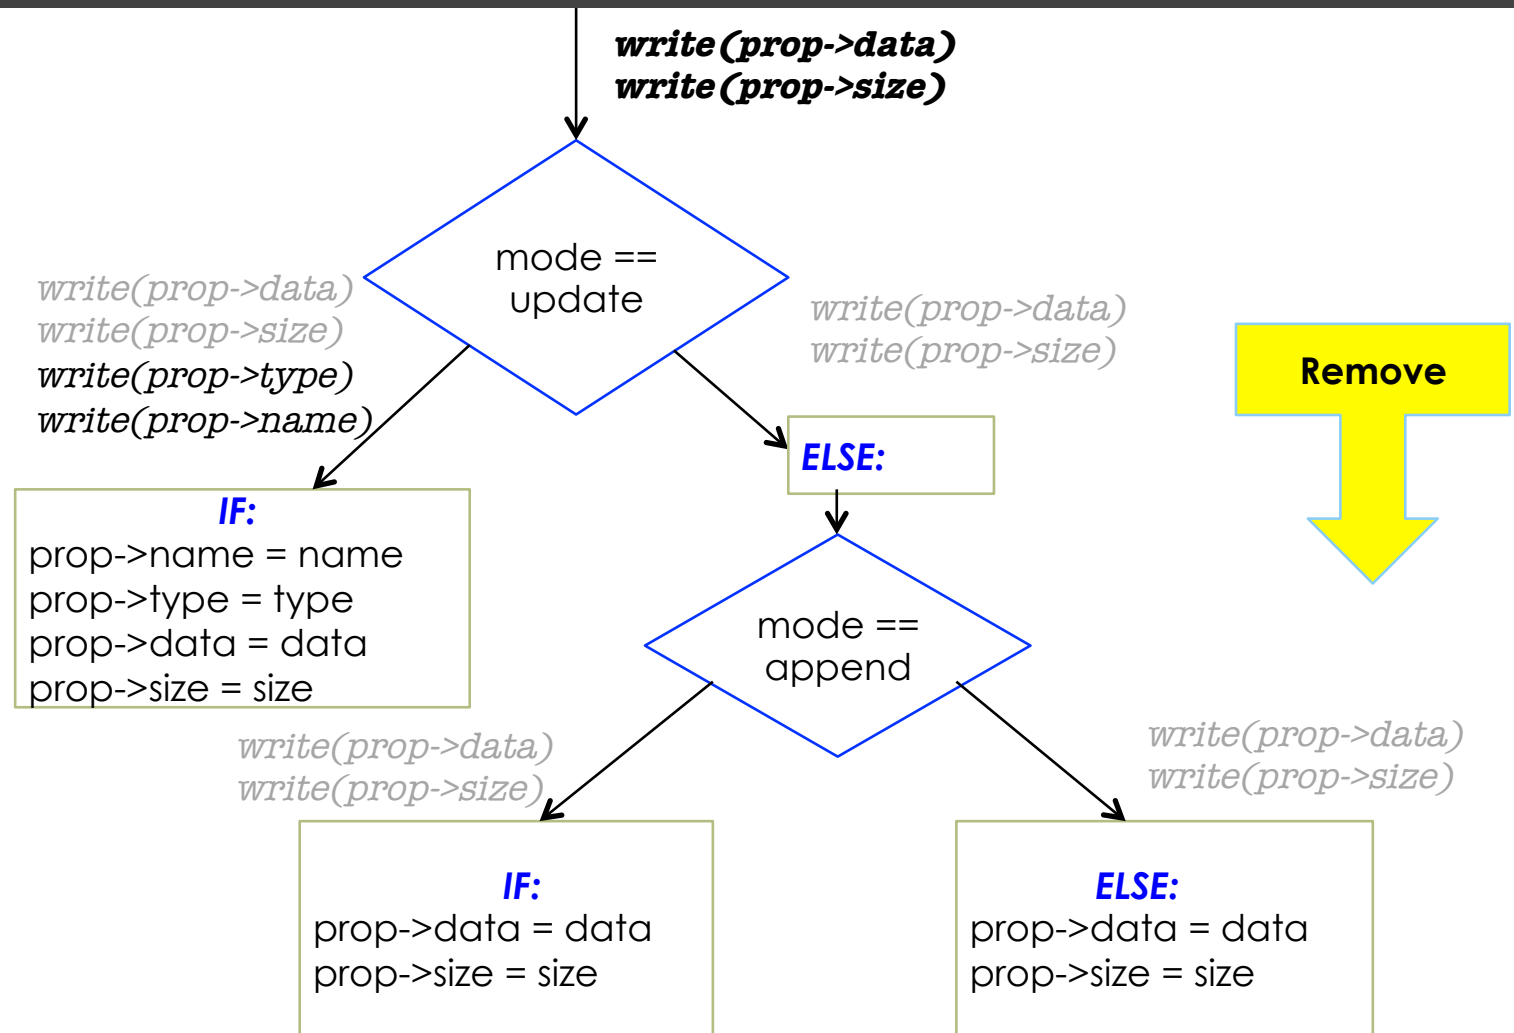
Non-redundant Hook Placement



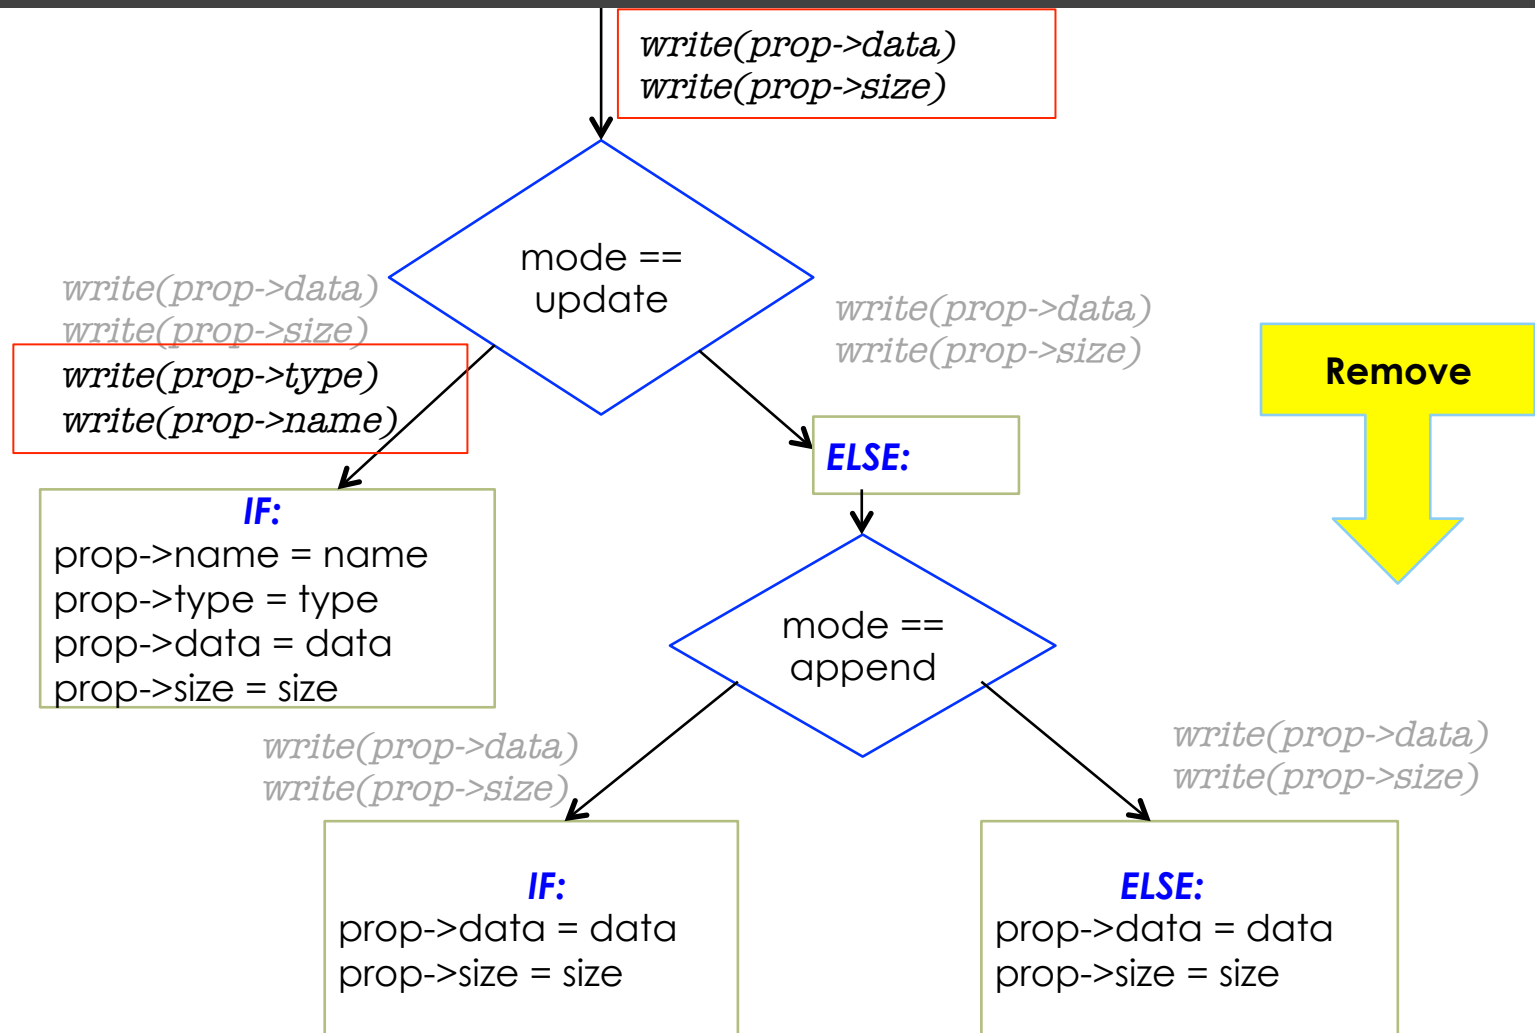
Non-redundant Hook Placement



Non-redundant Hook Placement



Non-redundant Hook Placement



Results

Program	X Server	postgres	pennmush	memcached
LOC	28k	49k	78k	9k
Total variables	7795	12350	24372	2350
Tainted variables	2975 (38%)	5100 (41%)	4188 (17%)	490 (20%)
Security sensitive variables	823 (10%)	402 (3%)	1573 (6%)	82 (3%)
Data Structures	404	278	311	41
Sensitive Data structures	61(15%)	30 (10%)	38 (12%)	7 (17%)
User-choice Operations	4760	5063	6485	996
Sensitive operations	1382 (29%)	1378 (27%)	1382 (21%)	203 (20%)
Hooks	532 (11%)	579 (11%)	714 (11%)	56 (5%)

Results

Program	X Server	postgres	pennmush	memcached
LOC	28k	49k	78k	9k
Total variables	7795	12350	24372	2350
Tainted variables	2975 (38%)	5100 (41%)	4168 (17%)	490 (20%)
Security sensitive variables	823 (10%)	402 (3%)	1573 (6%)	82 (3%)
Data Structures	404	278	311	41
Sensitive Data structures	61(15%)	30 (10%)	38 (12%)	7 (17%)
User-choice Operations	4760	5063	6485	996
Sensitive operations	1382 (29%)	1378 (27%)	1382 (21%)	203 (20%)
Hooks	532 (11%)	579 (11%)	714 (11%)	56 (5%)

Results

Program	X Server	postgres	pennmush	memcached
LOC	28k	49k	78k	9k
<i>Total variables</i>	7795	12350	24372	2350
Tainted variables	2975 (38%)	5100 (41%)	4168 (17%)	490 (20%)
Security sensitive variables	823 (10%)	402 (3%)	1573 (6%)	82 (3%)
Data Structures	404	278	311	41
Sensitive Data structures	61(15%)	30 (10%)	38 (12%)	7 (17%)
User-choice Operations	4760	5063	6485	996
Sensitive operations	1382 (29%)	1378 (27%)	1382 (21%)	203 (20%)
Hooks	532 (11%)	579 (11%)	714 (11%)	56 (5%)

Results

Program	X Server	postgres	pennmush	memcached
LOC	28k	49k	78k	9k
<i>Total variables</i>	7795	12350	24372	2350
Tainted variables	2575 (33%)	5100 (41%)	4168 (17%)	490 (20%)
Security sensitive variables	823 (10%)	402 (3%)	1573 (6%)	82 (3%)
Data Structures	311	311	311	41
Sensitive Data structures	311 (10%)	311 (10%)	38 (12%)	7 (17%)
User-choice Operations	4750	5063	6485	996
Sensitive operations	1332 (28%)	1378 (27%)	1382 (21%)	203 (20%)
Hooks	532 (11%)	579 (11%)	714 (11%)	56 (5%)

90+% effort reduction

Limitations

- Alias Analysis:
 - Cannot prove *minimality* without *complete* alias analysis.
 - Cannot prove *completeness* without assuming *sound* alias analysis.
- Implicit flows:
 - Typically cause an unwieldy number of false positives in static analysis.

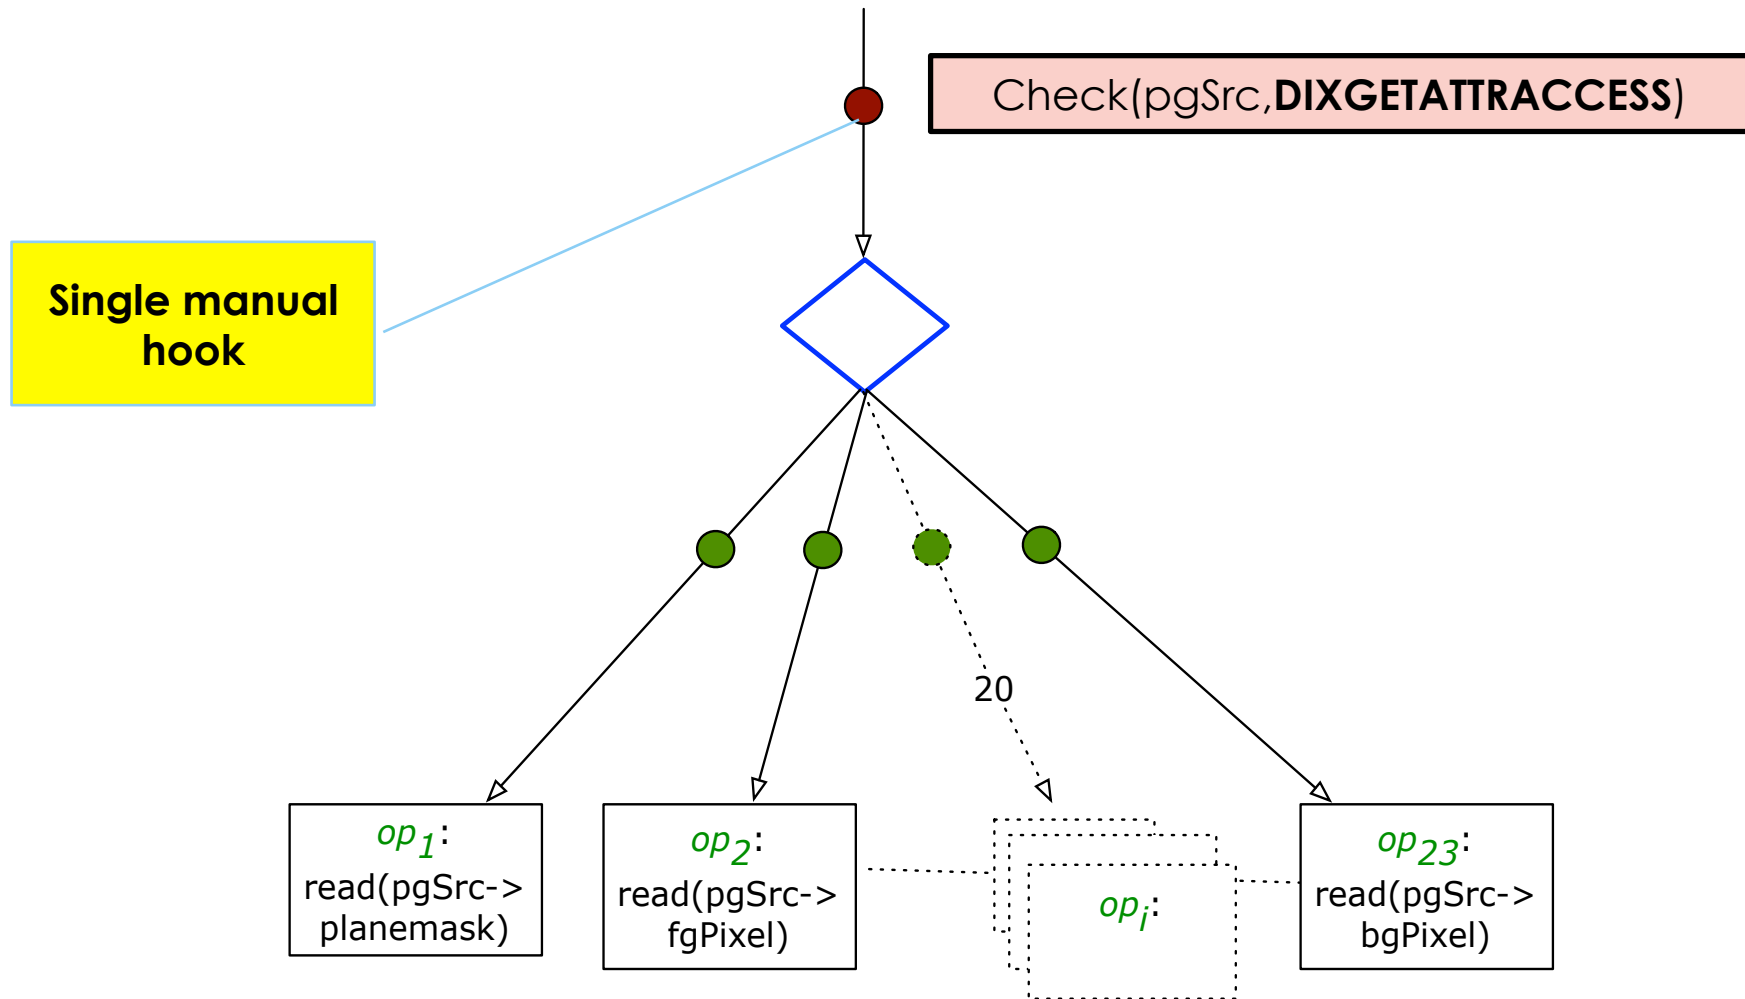
Comparison to Manual Hooks

Automated hook placement finer-grained than manual placement

- **X Server (version 1.9 with XACE hooks):**
 - Manual: 207 hooks
 - Automated: 532 hooks
- **Postgres (version 9.0 with sepgsql hooks):**
 - Manual: 370 hooks
 - Automated: 579 hooks

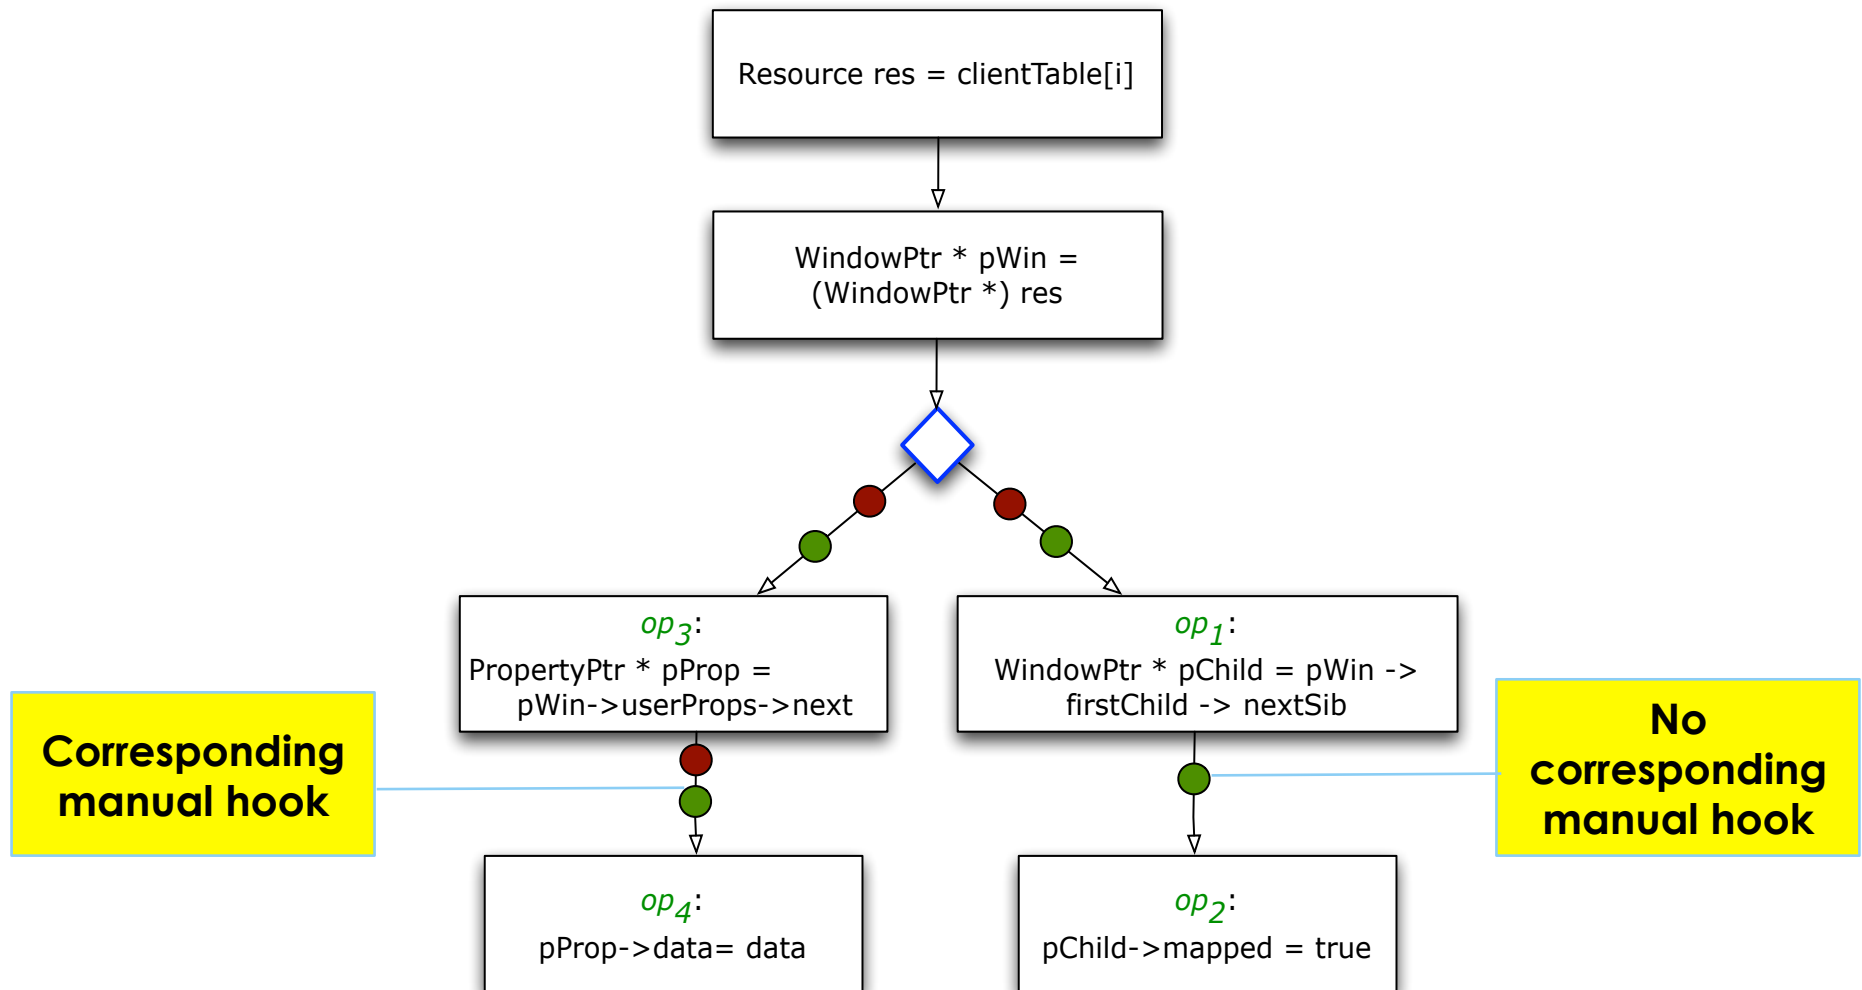
Claim: Placements must be *minimal* with respect to the *expected* authorization *policy*.

Comparison to Manual Hooks



CopyGC @ gc.c (X Server 1.1.3)

Comparison to Manual Hooks



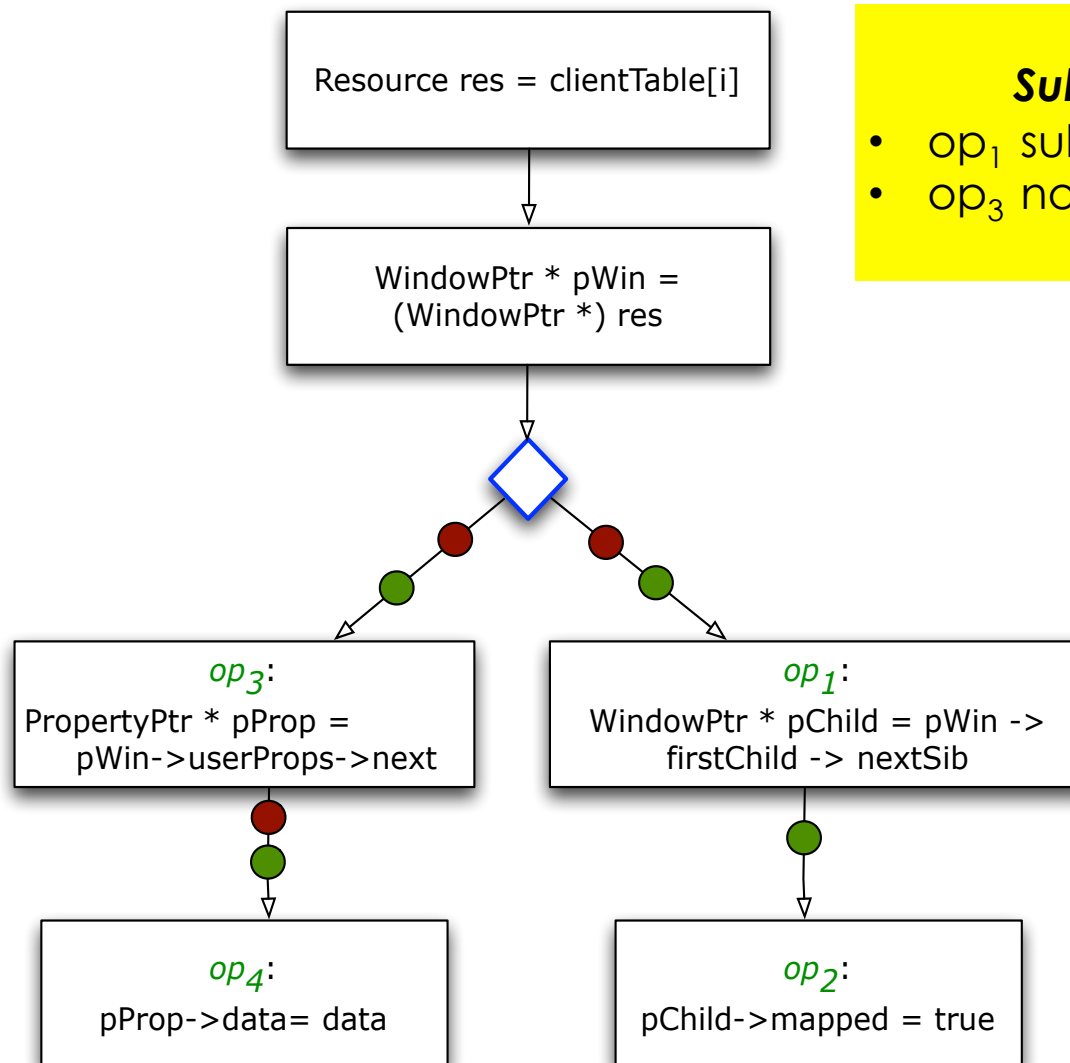
Authorization Constraints

- $Allowed(o)$: Subset of subjects in U that are allowed to perform operation o .
- **Constraint I:**
 - $Allowed(o1) = Allowed(o2)$, then **$o1$ equals $o2$**
- **Constraint II:**
 - $Allowed(o1) \subset Allowed(o2)$, then **$o1$ subsumes $o2$**

Authorization Constraints

- $Allowed(o)$: Subset of subjects in U that are allowed to perform operation o .
 - Suppose operation $o1$ control-flow dominates operation $o2$
- **Constraint I:**
 - $Allowed(o1) = Allowed(o2)$, then **$o1$ equals $o2$**
- **Constraint II:**
 - $Allowed(o1) \subset Allowed(o2)$, then **$o1$ subsumes $o2$**

Constraints: *Subsumption*



Subsumption:

- op_1 subsumes op_2
- op_3 not subsumes op_4

Authorization Constraints

- $Allowed(o)$: Subset of subjects in U that are allowed to perform operation o .
- Invariant I:
 - $Allowed(o1) = Allowed(o2)$, then **$o1$ equals $o2$**
- Invariant II:
 - $Allowed(o1) \subset Allowed(o2)$, then **$o1$ subsumes $o2$**

- **Access Control Policy is *not* available.**
- **Access Control Policy is not in terms of *code level operations***

Challenges

- *How to generate a placement that minimizes hooks w.r.t authorization constraints?*
- *How do we get authorization constraints?*

Challenges

- *How to generate a placement that minimizes hooks w.r.t authorization constraints?*
 - Use authorization constraints to eliminate “redundant” hooks
- *How do we get authorization constraints?*

Challenges

- *How to generate a placement that minimizes hooks w.r.t. authorization constraints?*
 - Use authorization constraints to eliminate “redundant” hooks
- *How do we get authorization constraints?*
 - “Top-down”
 - Programmers propose placement and we compute authorization constraints
 - Programmers choose authorization constraints
 - “Bottom-up”
 - Start from any placement, such as computed default
 - Compute constraints relative to that placement
 - Could “select” a group of constraints that satisfy a high-level constraint automatically

Constraints and Placements

**Default
Placement**

*Can enforce any access
control policy.*

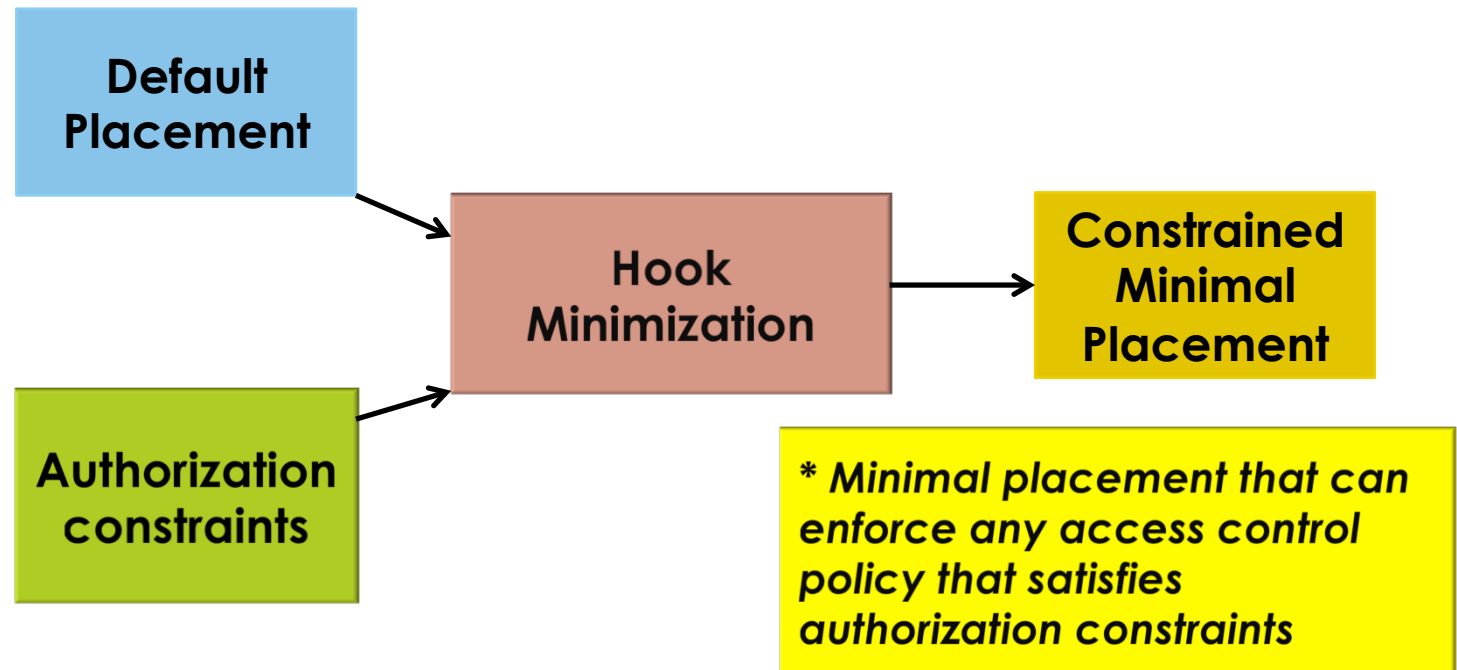
**Authorization
constraints**

*Equivalence and subsumption
relationships on operations*

**Candidate
Placement**

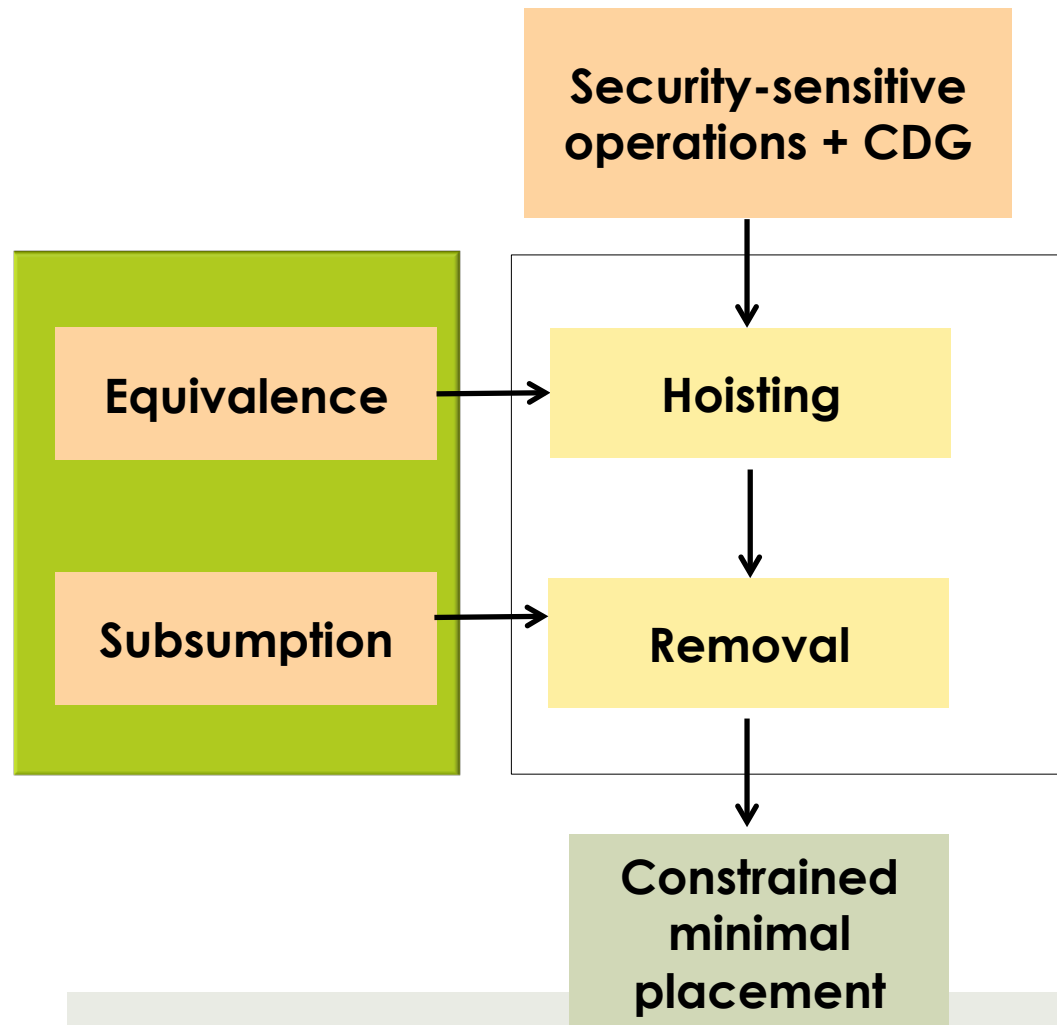
*Can enforce any access
control policy that satisfies
authorization constraints*

Compute Minimal Placement




Generating a minimal placement


- Equivalence (Q) + Subsumption (S)



Top-Down: Infer Authorization Constraints




Bottom-Up: Infer Authorization Constraints



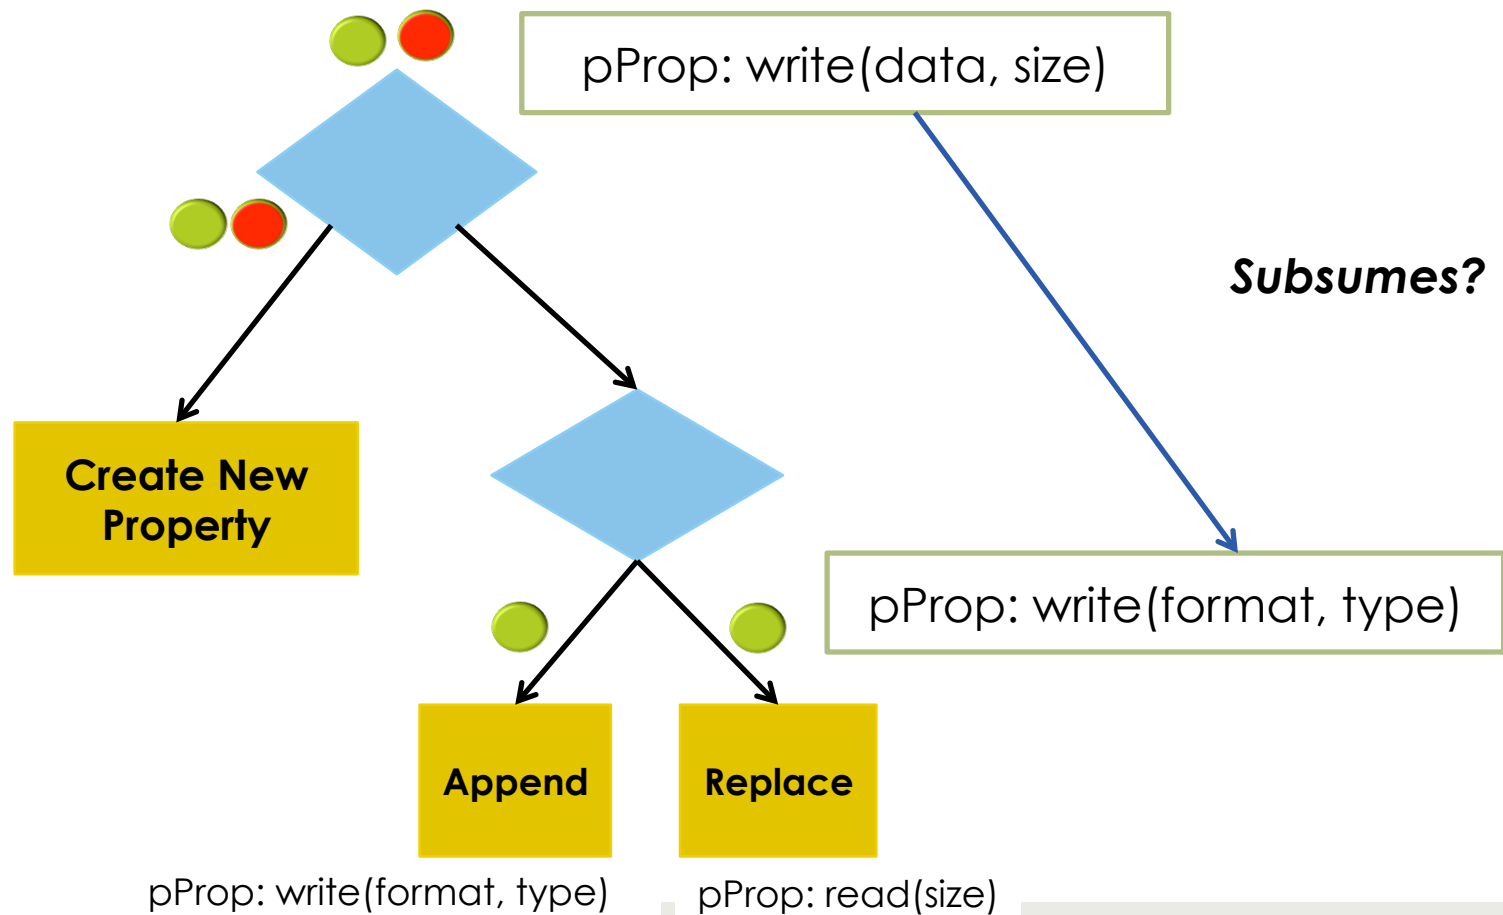
Show additional choices

- ▣ Given a placement,
 - ▣ What additional hoisting can be done?
 - ▣ Sets of control statement hooks with a common control statement.
 - ▣ What additional removal can be done?
 - ▣ Hook that have at least one dominating hook.

Additional choices



Subsumption choices



Helping programmers infer constraints

Top Down:

- Programmers guess at placement
- We show au implied.

Constraint Selectors:


At each hoisting or removal point apply rules to decide whether the operations are equivalent or subsuming

Bottom Up:

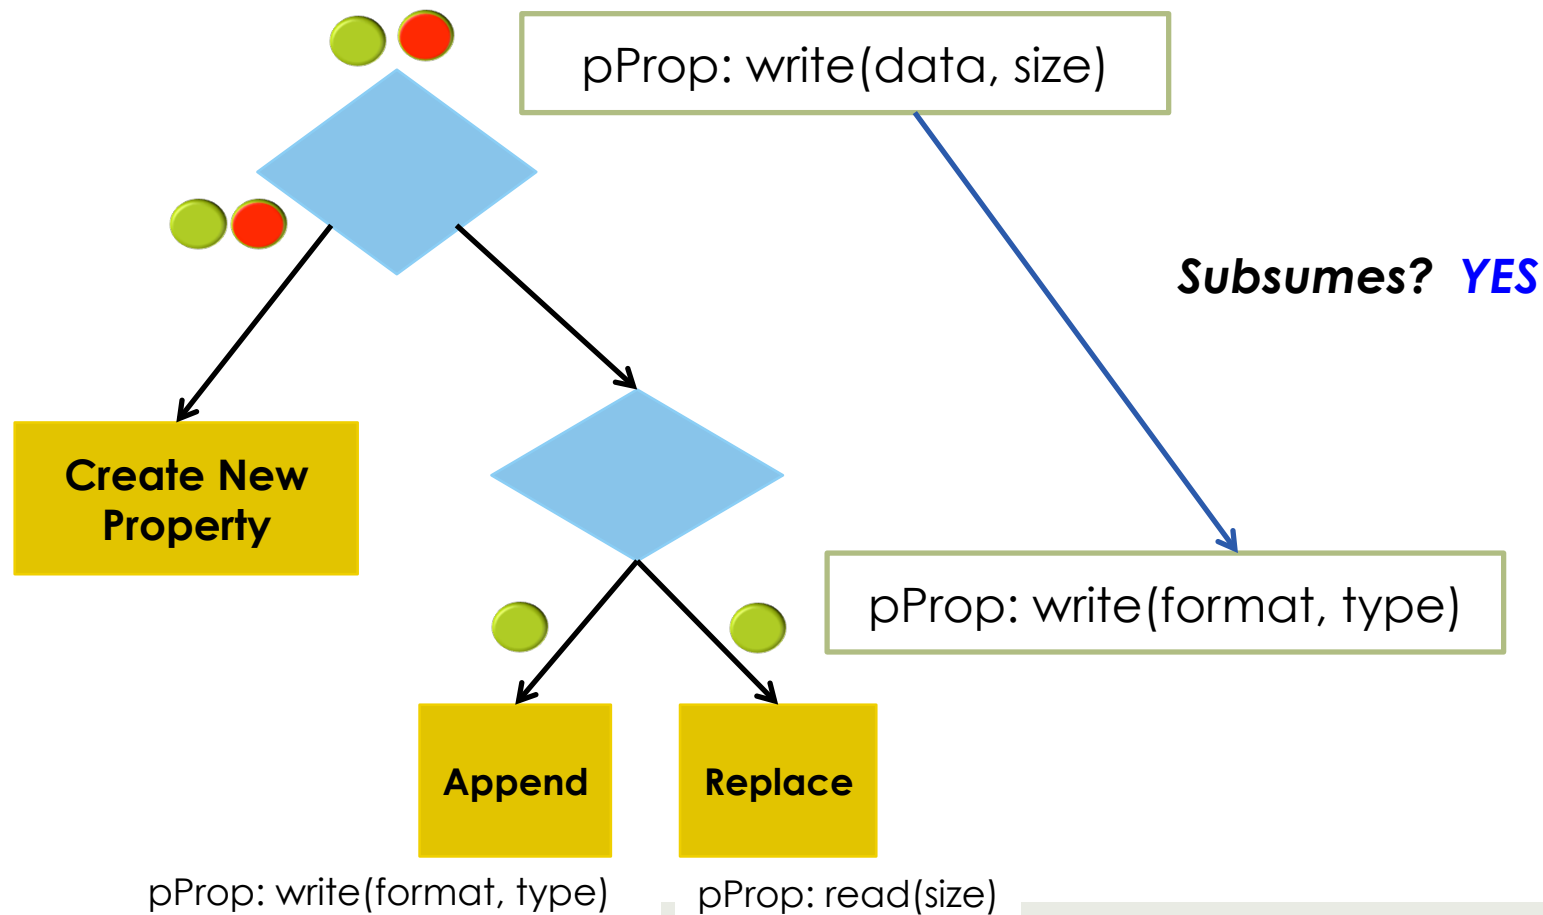
placement (say
w what hoists and
ill happen.

- Advise the programmer about constraints that may be useful.

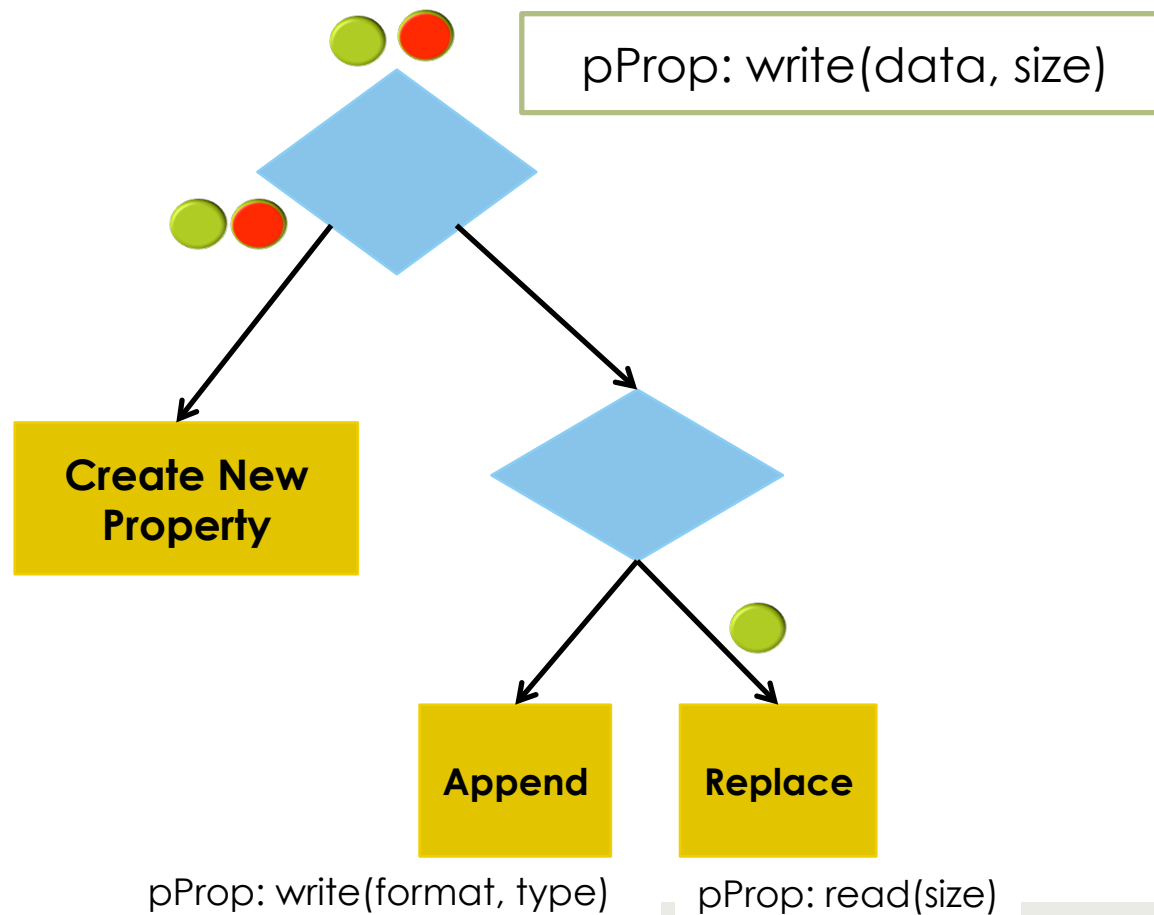
MLS Selector Example



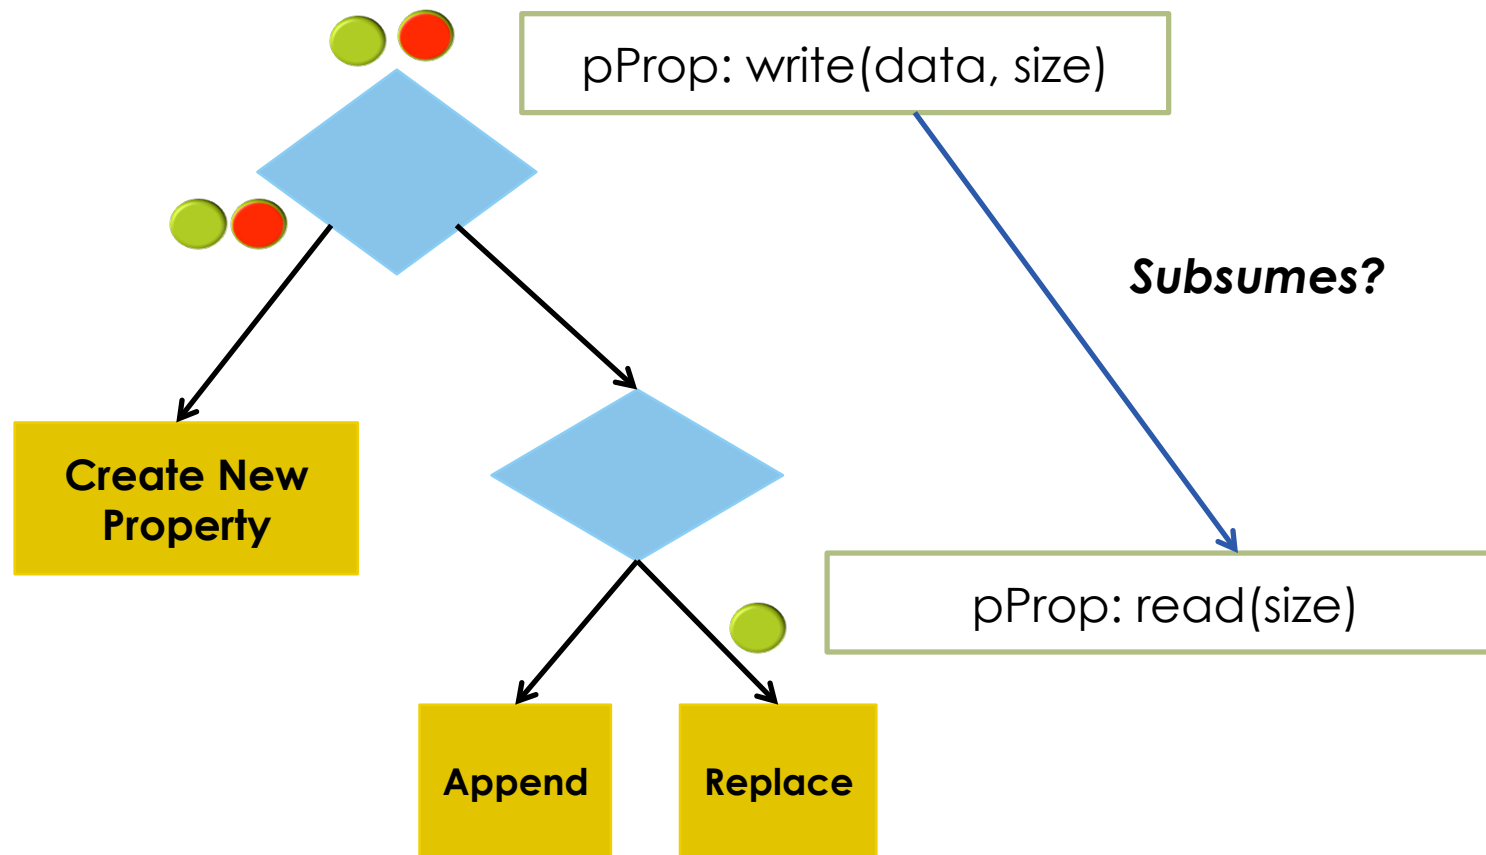
MLS Selector Example



MLS Selector Example



MLS Selector Example



Experiments

- 4 programs:
 - X Server 1.13 ([manual](#))
 - Postgres 9.1.9 ([manual](#))
 - Linux Kernel VFS 2.6.38.8 ([manual](#))
 - memcached
- 4 selectors:
 - DEF
 - FIL-MLS
 - FIL-RUN
 - FIL-MLS + FIL-RUN

Experiments

- How many authorization constraints do programmers have to look at?
 - How do selectors help with that?
- How many hoisting and removal choices do programmers have to make?
 - How do selectors help with that?